


第一次成績分佈 人數



Max: 97 Avg: 55

**c**

1. Which keyword can be used to check whether the result of a correlated nested query is empty or not? (5%)  
(a) DISTINCT (b) ALL (c) EXISTS (d) IN.

**a d**

2. Which mechanisms can be used to enforce business rules? (5%)  
(a) CREATE ASSERTION (b) ALTER TABLE (c) CREATE TABLE (d) CREATE TRIGGER.

3. Define the following terms: (a) data model (b) superkey (c) naïve or parametric end-users (d) internal schema. (20%)
4. In case of integrity violation in update operations on relations, please describe what actions can be taken? (10%)

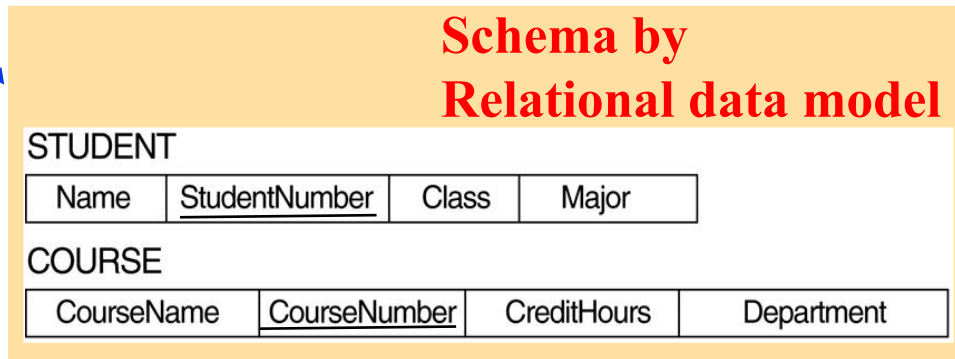


# Data Models

- **Data Model**

- A set of concepts to describe the *structure* of a database, and certain *constraints* that the database should obey.

- ER data model: {entity, relationship, attribute, key, ...}

- Relational data model: {relation, tuple, attribute, primary key, ...}



# Key Constraints

- **Superkey of R:**
  - A set of attributes **SK** of R such that **no two tuples in any valid relation instance  $r(R)$  will have the same value for SK.**
  - For any distinct tuples  $t_1$  and  $t_2$  in  $r(R)$ ,  $t_1[SK] \neq t_2[SK]$ .  
SK={LicenseNumber}?, SK={Make}?, SK={LicenseNumber, Make}?
- **Key of R:**
  - A "**minimal**" superkey;
  - A superkey K such that removal of any attribute from K results in a set of attributes that is not a superkey.
    - Key1 = {LicenseNumber}, Key2 = {EngineSerialNumber}  
are two keys of relation CAR; Key = {Dnumber, Dlocation}
    - SK1={EngineSerialNumber, Make}, SK2={LicenseNumber, Make, Year}  
are superkeys but **not** keys

CAR	<u>LicenseNumber</u>	EngineSerialNumber	Make	Model	Year
t1 →	Texas ABC-739	A69352	Ford	Mustang	96
t2 →	Florida TVP-347	B43696	Oldsmobile	Cutlass	99
	New York MPO-22	X83554	Oldsmobile	Delta	95
	California 432-TFY	C43742	Mercedes	190-D	93
	California RSK-629	Y82935	Toyota	Camry	98
	Texas RSK-629	U028365	Jaguar	XJS	98

DEPT\_LOCATIONS

<u>Dnumber</u>	<u>Dlocation</u>
1	Houston
4	Stafford
5	Bellaire
5	Sugarland
5	Houston

# Categories of End-users-1

- **Naïve or Parametric**


- They use previously well-defined functions in the form of “canned transactions” against the database.
- Examples are *bank-tellers* or *reservation clerks* who do this activity for an entire shift of operations.



# Three-Schema Architecture

- Defines DBMS schemas at *three levels*:
  - ✓ **External schemas**  
at the external level to describe the various user views. Usually uses the same data model as the conceptual level.
  - ✓ **Conceptual schema**  
at the conceptual level to describe the structure and constraints for the *whole* database for a community of users. Uses a *conceptual* or an *implementation* data model.
  - ✓ **Internal schema**  
at the internal level to describe physical storage structures and access paths. Typically uses a *physical* data model.

DBA  
defines DB  
schema



# Update Operations on Relations

- In case of integrity violation, several actions can be taken:
  - Cancel the operation that causes the violation (**REJECT** option)
  - Perform the operation but inform the user of the violation
  - Trigger additional updates so the violation is corrected (**CASCADE** option, **SET NULL** option, **SET DEFAULT** option)
  - Execute a user-specified error-correction routine

EMPLOYEE (PK)						(FK)		(FK)	
Fname	Minit	Lname	<u>Ssn</u>	Bdate	Address	Sex	Salary	Super_ssn	Dno
John	B	Smith	123456789	1965-01-09	731 Fondren, Houston, TX	M	30000	333445555	5
Franklin	T	Wong	333445555	1955-12-08	638 Voss, Houston, TX	M	40000	888665555	5
Alicia	J	Zelaya	999887777	1968-01-19	3321 Castle, Spring, TX	F	25000	987654321	4
Jennifer	S	Wallace	987654321	1941-06-20	291 Berry, Bellaire, TX	F	43000	888665555	4
Ramesh	K	Narayan	666884444	1962-09-15	975 Fire Oak, Humble, TX	M	38000	333445555	5
Joyce	A	English	453453453	1972-07-31	5631 Rice, Houston, TX	F	25000	333445555	5
Ahmad	V	Jabbar	987987987	1969-03-29	980 Dallas, Houston, TX	M	25000	987654321	4
James	E	Borg	888665555	1937-11-10	450 Stone, Houston, TX	M	55000	NULL	1

Delete

Modify

DEPARTMENT (PK)		(FK)	
Dname	<u>Dnumber</u>	Mgr_ssn	Mgr_start_date
Research	5	333445555	1988-05-22
Administration	4	987654321	1995-01-01
Headquarters	1	888665555	1981-06-19

WORKS\_ON

<u>Essn</u>	<u>Pno</u>	Hours
123456789	1	32.5
123456789	2	7.5
666884444	3	40.0
453453453	1	20.0
453453453	2	20.0
333445555	2	10.0
333445555	3	10.0

(FK)

(FK)



5. Please describe when not to use a DBMS. (10%)
6. SQL is a non-procedural language and C is a procedural language. Please describe their differences between non-procedural and procedural languages. (10%)

# When not to Use a DBMS

- **Main inhibitors (costs) of using a DBMS:**

- High initial investment and possible need for additional hardware.
- Overhead for providing generality, security, concurrency control, recovery, and integrity functions.

- **When a DBMS may be unnecessary:**

- If the database and applications are simple, well defined, and not expected to change.
- If there are stringent real-time requirements that may not be met because of DBMS overhead.
- If access to data by multiple users is not required.

- **When no DBMS may suffice:**

- If the database system is not able to handle the complexity of data because of modeling limitations
- If the database users need special operations not supported by the DBMS.



# DBMS Languages

- **High Level or Non-procedural Languages:**
  - Also called *declarative* languages.
  - e.g., SQL, are *set-oriented* and specify *what* data to retrieve than how to retrieve.
- **Low Level or Procedural Languages:**
  - *record-at-a-time*; they specify *how* to retrieve data and include constructs such as looping.

```
SELECT  SNO, NAME  
FROM    STUDENT  
WHERE   SEX = M
```

```
for (i = 0, TotalStuNum, i++)  
    if (Student[i].Sex == M)  
        print(Student[i].Sno, Student[i].Name)
```

- (a) Retrieve the birthdate and address of the employee whose name is 'John B. Smith'.
- (b) For each employee, retrieve the employee's name, and the name of his or her immediate supervisor.
- (c) Retrieve the names of employees who have no dependents.
- (d) Retrieve a list of employees and the projects each works in, ordered by the employee's department, and within each department ordered alphabetically by employee last name.
- (e) Show the effect of giving all employees who work on the 'ProductX' project a 10% raise.
- (f) For each department, retrieve the department number, the number of employees in the department, and their average salary.
- (g) Give all employees in department 5 a 10% raise in salary.
- (h) Delete the employees whose address is in Douliu, Yunlin.

# Simple SQL Queries

- Basic SQL queries correspond to using the SELECT, PROJECT, and JOIN operations of the relational algebra
- Example of a simple query on *one* relation
- Query 0: Retrieve the birthdate and address of the employee whose name is 'John B. Smith'.

**Q0: SELECT BDATE, ADDRESS  
FROM EMPLOYEE  
WHERE FNAME='John' AND MINIT='B' AND LNAME='Smith'**

- Similar to a SELECT-PROJECT pair of relational algebra operations; the SELECT-clause specifies the *projection attributes* and the WHERE-clause specifies the *selection condition*
- However, the result of the query *may contain* duplicate tuples

BDATE	ADDRESS
1965-01-09	731 Fondren, Houston, TX

EMPLOYEE									
FNAME	MINIT	LNAME	<u>SSN</u>	BDATE	ADDRESS	SEX	SALARY	SUPERSSN	DNO



# ALIASES

- Some queries need to refer to the same relation twice
- In this case, *aliases* are given to the relation name
- Query 8: For each employee, retrieve the employee's name, and the name of his or her immediate supervisor.

**Q8:**    **SELECT**    E.FNAME, E.LNAME, S.FNAME, S.LNAME  
         **FROM**     EMPLOYEE E S  
         **WHERE**    E.SUPERSSN=S.SSN

- In Q8, the alternate relation names E and S are called *aliases* or *tuple variables* for the EMPLOYEE relation
- We can think of E and S as two *different copies* of EMPLOYEE; E represents employees in role of *supervisees* and S represents employees in role of *supervisors*
- Aliasing can also be used in any SQL query for convenience  
Can also use the AS keyword to specify aliases

**Q8:**    **SELECT**    E.FNAME, E.LNAME, S.FNAME, S.LNAME  
         **FROM**     EMPLOYEE AS E, EMPLOYEE AS S  
         **WHERE**    E.SUPERSSN=S.SSN

**E** ↓

EMPLOYEE	FNAME	MINIT	LNAME	SSN	BDATE	ADDRESS	SEX	SALARY	SUPERSSN	DNO
	John	B	Smith	123456789	1965-01-09	731 Fondren, Houston, TX	M	30000	333445555	5
	Franklin	T	Wong	333445555	1955-12-08	638 Voss, Houston, TX	M	40000	888665555	5

**S** ↓


EMPLOYEE	FNAME	MINIT	LNAME	SSN	BDATE	ADDRESS	SEX	SALARY	SUPERSSN	DNO
	John	B	Smith	123456789	1965-01-09	731 Fondren, Houston, TX	M	30000	333445555	5
	Franklin	T	Wong	333445555	1955-12-08	638 Voss, Houston, TX	M	40000	888665555	5

# THE EXISTS FUNCTION

- Query 6: Retrieve the names of employees who have no dependents.

**Q6: SELECT    FNAME, LNAME  
         FROM    EMPLOYEE  
         WHERE   NOT EXISTS ( SELECT   \*  
                                 FROM    DEPENDENT  
                                 WHERE   SSN=ESSN)**

- In Q6, the correlated nested query retrieves all DEPENDENT tuples related to an EMPLOYEE tuple. If *none exist*, the EMPLOYEE tuple is selected
- EXISTS is necessary for the expressive power of SQL



# ORDER BY

- The **ORDER BY** clause is used to sort the tuples in a query result based on the values of some attribute(s)

Query 15: Retrieve a list of employees and the projects each works in, ordered by the employee's department, and within each department ordered alphabetically by employee last name.

```
Q15: SELECT      DNAME, LNAME, FNAME, PNAME  
      FROM        DEPARTMENT, EMPLOYEE, WORKS_ON, PROJECT  
      WHERE       DNUMBER=DNO AND SSN=ESSN AND PNO=PNUMBER  
      ORDER BY    DNAME, LNAME
```

- The default order is in ascending order of values
- Keyword **DESC** if we want a descending order;
- Keyword **ASC** can be used to explicitly specify ascending order, even though it is the default



# ARITHMETIC OPERATIONS

- The standard arithmetic operators '+', '-', '\*', and '/' (for addition, subtraction, multiplication, and division, respectively) can be applied to numeric values in an SQL query result
- Query 13: Show the effect of giving all employees who work on the 'ProductX' project a 10% raise.

**Q13: SELECT FNAME, LNAME, 1.1\*SALARY  
FROM EMPLOYEE, WORKS\_ON, PROJECT  
WHERE SSN=ESSN AND PNO=PNUMBER AND  
PNAME='ProductX'**



EMPLOYEE	FNAME	MINIT	LNAME	SSN	BDATE	ADDRESS	SEX	SALARY	SUPERSSN	DNO
	John	B	Smith	123456789	1965-01-09	731 Fondren, Houston, TX	M	30000	333445555	5
	Franklin	T	Wong	333445555	1955-12-08	638 Voss, Houston, TX	M	40000	888665555	5
	Alicia	J	Zelaya	999887777	1968-07-19	3321 Castle, Spring, TX	F	25000	987654321	4
	Jennifer	S	Wallace	987654321	1941-06-20	291 Berry, Bellaire, TX	F	43000	888665555	4
	Ramesh	K	Narayan	666884444	1962-09-15	975 Fire Oak, Humble, TX	M	38000	333445555	5
	Joyce	A	English	453453453	1972-07-31	5631 Rice, Houston, TX	F	25000	333445555	5
	Ahmad	V	Jabbar	987987987	1969-03-29	980 Dallas, Houston, TX	M	25000	987654321	4
	James	E	Borg	888665555	1937-11-10	450 Stone, Houston, TX	M	55000	null	1

# GROUPING (cont.)

- Query 24: For each department, retrieve the department number, the number of employees in the department, and their average salary.

**Q24: SELECT            DNO, COUNT (\*), AVG (SALARY)  
                      FROM            EMPLOYEE  
                      GROUP BY     DNO**

- the EMPLOYEE tuples are divided into groups--each group having the same value for the grouping attribute DNO
- The COUNT and AVG functions are applied to each such group of tuples separately
- The SELECT-clause includes only the grouping attribute and the functions to be applied on each group of tuples

FNAME	MINIT	LNAME	SSN	• • •	SALARY	SUPERSSN	DNO
John	B	Smith	123456789	• • •	30000	333445555	5
Franklin		Wong	333445555		40000	888665555	5
Ramesh	K	Narayan	666884444		38000	333445555	5
Joyce	A	English	453453453		25000	333445555	5
Alicia	J	Zelaya	999887777		25000	987654321	4
Jennifer	S	Wallace	987654321		43000	888665555	4
Ahmad	V	Jabbar	987987987		25000	987654321	4
James	E	Bong	888665555		55000	null	1

Grouping EMPLOYEE tuples by the value of DNO.

DNO	COUNT (*)	AVG (SALARY)
5	4	33250
4	3	31000
1	1	55000

Result of Q24.

# UPDATE (cont.)

- Example: Give all employees in department 5 a 10% raise in salary.

**U6: UPDATE EMPLOYEE**  
**SET SALARY = SALARY \* 1.1**  
**WHERE DNO = 5**

- In this request, the modified SALARY value depends on the original SALARY value in each tuple
- The reference to the SALARY attribute on the right of = refers to the old SALARY value before modification
- The reference to the SALARY attribute on the left of = refers to the new SALARY value after modification

EMPLOYEE	FNAME	MINIT	LNAME	SSN	BDATE	ADDRESS	SEX	SALARY	SUPERSSN	DNO
	John	B	Smith	123456789	1965-01-09	731 Fondren, Houston, TX	M	30000	333445555	5
	Franklin	T	Wong	333445555	1955-12-08	638 Voss, Houston, TX	M	40000	888665555	5
	Alicia	J	Zelaya	999887777	1968-07-19	3321 Castle, Spring, TX	F	25000	987654321	4
	Jennifer	S	Wallace	987654321	1941-06-20	291 Berry, Bellaire, TX	F	43000	888665555	4
	Ramesh	K	Narayan	666884444	1962-09-15	975 Fire Oak, Humble, TX	M	38000	333445555	5
	Joyce	A	English	453453453	1972-07-31	5631 Rice, Houston, TX	F	25000	333445555	5
	Ahmad	V	Jabbar	987987987	1969-03-29	980 Dallas, Houston, TX	M	25000	987654321	4
	James	E	Borg	888665555	1937-11-10	450 Stone, Houston, TX	M	55000	null	1

# DELETE (cont.)

U4C: **DELETE FROM EMPLOYEE**  
**WHERE ADDRESS LIKE '%Douliu, Yunlin%'**

A missing WHERE-clause specifies that *all tuples* in the relation are to be deleted; the table then becomes an empty table

EMPLOYEE	FNAME	MINIT	LNAME	<u>SSN</u>	BDATE	ADDRESS	SEX	SALARY	SUPERSSN	DNO
	John	B	Smith	123456789	1965-01-09	731 Fondren, Houston, TX	M	30000	333445555	5
	Franklin	T	Wong	333445555	1955-12-08	638 Voss, Houston, TX	M	40000	888665555	5
	Alicia	J	Zelaya	999887777	1968-07-19	3321 Castle, Spring, TX	F	25000	987654321	4
	Jennifer	S	Wallace	987654321	1941-06-20	291 Berry, Bellaire, TX	F	43000	888665555	4
	Ramesh	K	Narayan	666884444	1962-09-15	975 Fire Oak, Humble, TX	M	38000	333445555	5
	Joyce	A	English	453453453	1972-07-31	5631 Rice, Houston, TX	F	25000	333445555	5
	Ahmad	V	Jabbar	987987987	1969-03-29	980 Dallas, Houston, TX	M	25000	987654321	4
	James	E	Borg	888665555	1937-11-10	450 Stone, Houston, TX	M	55000	null	1