

Advanced SQL Techniques

Part 1

Selecting the First Records

- Quite often, it's necessary to find the first records based on some criteria.
- For example, let's take the `car_portal` database; suppose you need to find the earliest advertisement for each `car_id` in the advertisement table.

```
car_portal=> SELECT advertisement_id, advertisement_date, adv.car_id, seller_account_id
car_portal-> FROM   car_portal_app.advertisement adv INNER JOIN (SELECT   car_id, min(advertisement_date) min_date
car_portal(>                                     FROM     car_portal_app.advertisement
car_portal(>                                     GROUP BY car_id ) first
car_portal->      ON adv.car_id=first.car_id AND
car_portal->      adv.advertisement_date = first.min_date;
```

- However, if the ordering logic is complex and cannot be implemented using the `min` function, this approach won't work.

- Although window functions can solve the problem, they aren't always convenient to use:

```
car_portal=> SELECT DISTINCT first_value(advertisement_id) OVER w AS advertisement_id,  
car_portal->      min(advertisement_date) OVER w AS advertisement_date,  
car_portal->      car_id,  
car_portal->      first_value(seller_account_id) OVER w AS seller_account_id  
car_portal-> FROM    car_portal_app.advertisement  
car_portal-> WINDOW w AS (PARTITION BY car_id ORDER BY advertisement_date);
```

- In the preceding code, `DISTINCT` is used to remove the duplicates that were grouped together in the previous example.

- PostgreSQL provides an explicit way of selecting the first record within each group.
- The `DISTINCT ON` construct is used for this.
- The syntax is as follows:
 - `SELECT DISTINCT ON (<expression_list>) <Select-List>`
 `...`
 `ORDER BY <order_by_list>`

- In the preceding code snippet, for each distinct combination of values of `expression_list`, only the first record will be returned by the `SELECT` statement.
- The `ORDER BY` clause is used to define a rule to determine which record is the first.
- The `expression_list` from the `DISTINCT ON` construct must be included in the `order_by_list` list.

```
car_portal=> SELECT    DISTINCT ON (car_id) advertisement_id, advertisement_date, car_id, seller_account_id
car_portal-> FROM      car_portal_app.advertisement
car_portal-> ORDER BY  car_id, advertisement_date;
```


Selecting a Data Sample

- Suppose the `car_portal` database is huge and there is a task to collect some statistics about the data.
- The statistics don't have to be exactly accurate, estimated values would be good enough.
- It's important to collect the data quickly.
- PostgreSQL provides a way to query a random fraction of a table's records.

- The syntax is as follows:

- ```
SELECT ...
FROM <table>
TABLESAMPLE <sampling_method> (<argument> [, ...])
[REPEATABLE (<seed>)]
```

- The `sampling_method` can be either `BERNOULLI` or `SYSTEM`—these are provided with a standard PostgreSQL installation.
- More sampling methods are available as extensions.
- Both the sampling methods take one argument, which is a probability that a row will be selected, as a percentage.

- The difference between the two methods is that `BERNOULLI` implies a full table scan and the algorithm will decide for each individual row whether it should be selected or not.
- The `SYSTEM` method does the same but on the level of blocks of rows.
  - The latter is significantly faster, but the results are less accurate because of the different numbers of rows in different blocks.

- The `REPEATABLE` keyword is used to specify the seed value used by the random functions in the sampling algorithm.
- The results of queries with the same seed value is the same, as long as the data in the table doesn't change.

- Sampling happens at the moment the data from a table is read, before any filtering, grouping, or other operations are performed.

- To illustrate sampling, let's increase the number of records in the advertisement table by executing the following query:

```
car_portal=> INSERT INTO car_portal_app.advertisement (advertisement_date, car_id, seller_account_id)
car_portal-> SELECT advertisement_date, car_id, seller_account_id
car_portal-> FROM car_portal_app.advertisement;
INSERT 0 784
```

- This simply duplicates the number of records in the table.
- If this is executed many times, the table will be pretty big.
- For the following example, that query was executed 12 more times.
- If we try to select the total count of records, it takes a fair bit of time:

```
car_portal=> \timing
Timing is on.
car_portal=> SELECT count(*) FROM car_portal_app.advertisement;
 count

 6422528
(1 row)

Time: 172.417 ms
```



- Now let's try sampling:

```
car_portal=> SELECT count(*) * 100
car_portal-> FROM car_portal_app.advertisement
car_portal-> TABLESAMPLE SYSTEM (1);
?column?

 6452700
(1 row)

Time: 7.763 ms
```

- The second query is much faster; however, the result is not very accurate.
- We multiply the count function by 100 because only one percent of the rows is selected.

- If this query is performed several times, the results will always be different and, sometimes, can be quite different from the real count of rows in the table.
- To improve quality, the fraction of the rows that are selected can be increased:

```
car_portal=> SELECT count(*) * 10
car_portal-> FROM car_portal_app.advertisement
car_portal-> TABLESAMPLE SYSTEM (10);
?column?

 6604990
(1 row)

Time: 62.669 ms
```

- The time increased proportionally.

# Set-Returning Functions

- Functions in PostgreSQL can return not only single values, but also relations.
- Such functions are called **set-returning functions**.

- This is a typical task for every SQL developer: generate a sequence of numbers, each in a separate record.
- This sequence or relation can have many use cases.
- For example, suppose you need to select data from the `car_portal` database to count the number of cars for each year of manufacture, from 2010 till 2015, and you need to show zero for the years where no cars exist in the system at all.
- A simple SELECT query from the only car table cannot be used to implement this logic.
- It isn't possible to count records that are not present in the table data.

- That's why it would be useful if there was a table with the numbers from 2010 to 2015.
- Then it could be outer-joined to the results of the query.

- You could create this table in advance, but that's not a very good idea, because the number of records for such a table wouldn't be known in advance, and if you created a very big table, it would be a waste of disk space in most cases.

- There is a function, `generate_series`, that serves exactly this purpose: it returns a set of integers with the given start and stop values.
- The query for the example would be as follows:

```
car_portal=> SELECT years.manufacture_year, count(car_id)
car_portal-> FROM generate_series(2010, 2015) as years (manufacture_year) LEFT JOIN car_portal_app.car
car_portal-> ON car.manufacture_year = years.manufacture_year
car_portal-> GROUP BY years.manufacture_year
car_portal-> ORDER BY 1;
 manufacture_year | count
-----+-----
 2010 | 11
 2011 | 12
 2012 | 12
 2013 | 21
 2014 | 17
 2015 | 0
(6 rows)
```

- It's possible to specify a step when calling `generate_series`:

```
car_portal=> SELECT *
car_portal-> FROM generate_series(5, 11, 3);
generate_series

 5
 8
 11
(3 rows)
```



- The `generate_series` function can also return a set of timestamp values:

```
car_portal=> SELECT *
car_portal-> FROM generate_series('2015-01-01'::date, '2015-01-31'::date, interval '7 days');
 generate_series

2015-01-01 00:00:00+08
2015-01-08 00:00:00+08
2015-01-15 00:00:00+08
2015-01-22 00:00:00+08
2015-01-29 00:00:00+08
(5 rows)
```

- There are a couple of other set-returning functions designed to work with arrays:
  - `generate_subscripts`: This generates numbers that can be used to reference the elements in a given array for the given dimension(s). This function is useful for enumerating the elements of an array in a SQL statement.
  - `unnest`: This transforms a given array into a set or rows where each record corresponds to an element of the array.

- Set-returning functions are also called **table functions**.

- Table functions can return sets of rows of a predefined type, like `generate_series` returns a set of `int` or `bigint` (depending on the type of input argument).
- Moreover, they can return a set of abstract type records.
- This allows a function to return different numbers of columns depending on their arguments.
- SQL requires that the row structure of all input relations is defined so that the structure of the result set will also be defined.
- That's why table functions that return sets of records, when used in a query, can be followed by a definition of their row structure:
  - `function_name (<arguments>) [AS] alias (column_name column_type [, ...])`

- The output of several set-returning functions can be combined together as if they were joined on the position of each row.
- The `ROWS FROM` syntax construct is used for this, as follows:
  - `ROWS FROM (function_call [,...]) [[AS] alias (column_name [, ...])]`
- The preceding construct will return a relation, therefore, it can be used in a `FROM` clause just like any other relation.
- The number of rows is equal to the largest output of the functions.
- Each column corresponds to the respective function in the `ROWS FROM` clause.
- If a function returns fewer rows than other functions, the missing values will be set to `NULL`.

```
car_portal=> SELECT foo.a, foo.b
car_portal-> FROM ROWS FROM (generate_series(1,3), generate_series(1,7,2)) AS foo(a, b);
 a | b
---+---
 1 | 1
 2 | 3
 3 | 5
 | 7
(4 rows)
```