

The PostgreSQL PL/pgSQL control statements

Declaration statements

- A **declaration statement** allows the developer to define variables and constants.
- The general syntax of a variable declaration is as follows:

```
name [ CONSTANT ] type [ COLLATE collation_name ] [ NOT NULL ] [ {  
DEFAULT | := | = } expression ];
```

- The PostgreSQL PL/pgSQL function body is composed of nested blocks, with an optional declaration section and a label.
- Variables are declared in the `DECLARE` section of the block, as follows:

```
[ <<label>> ]  
[ DECLARE  
    declarations ]  
BEGIN  
    statements  
END [ label ];
```

- To understand the function block, let's take a look at the following code, which defines the `factorial` function in a recursive manner:

```
car_portal=# CREATE OR REPLACE FUNCTION factorial(INTEGER ) RETURNS INTEGER AS $$
car_portal$# BEGIN
car_portal$#     IF $1 IS NULL OR $1 < 0 THEN RAISE NOTICE 'Invalid Number';
car_portal$#         RETURN NULL;
car_portal$#     ELIF $1 = 1 THEN
car_portal$#         RETURN 1;
car_portal$#     ELSE
car_portal$#         RETURN factorial($1 - 1) * $1;
car_portal$#     END IF;
car_portal$# END;
car_portal$# $$ LANGUAGE 'plpgsql';
CREATE FUNCTION
car_portal=#
car_portal=# select factorial(5);
factorial
-----
        120
(1 row)
```

- The block defines the variable scope; in our example, the scope of the argument variable, \$1, is the whole function.
- Also, as shown in the example, there is no declaration section.
- To understand the scope between different code blocks, let's write the factorial function in a slightly different manner, as follows:

```

car_portal=# CREATE OR REPLACE FUNCTION factorial(INTEGER ) RETURNS INTEGER AS $$
car_portal$#     DECLARE
car_portal$#         fact ALIAS FOR $1;
car_portal$#     BEGIN
car_portal$#         IF fact IS NULL OR fact < 0 THEN
car_portal$#             RAISE NOTICE 'Invalid Number';
car_portal$#             RETURN NULL;
car_portal$#         ELSIF fact = 1 THEN
car_portal$#             RETURN 1;
car_portal$#         END IF;
car_portal$#     DECLARE
car_portal$#         result INT;
car_portal$#     BEGIN
car_portal$#         result = factorial(fact - 1) * fact;
car_portal$#         RETURN result;
car_portal$#     END;
car_portal$# END;
car_portal$# $$ LANGUAGE 'plpgsql';
CREATE FUNCTION
car_portal=#
car_portal=# select factorial(5);
 factorial
-----
         120
(1 row)

```

- The preceding function is composed of two blocks; the `fact` variable is an alias for the first argument.
- In the sub-block, the `result` variable is declared with a type integer.
- Since the `fact` variable is defined in the upper block, it can also be used in the sub-block.
- The `result` variable can only be used in the sub-block.

Assignment statements

- An **assignment statement** is used to assign a value to a variable.
- Constants are assigned values at the time of declaration.
- The variable can be assigned an atomic value or a complex value, such as a record.
- Also, it can be assigned a literal value or a result of query execution.

- The assignment `:=` and `=`, operators are used to assign an expression to a variable, as follows:
`variable { := | = } expression;`
- For the variable names, you should choose names that do not conflict with the column names.
- This is important when writing parameterized SQL statements.
- The `=` operator is used in SQL for equality comparisons; it is preferable to use the `:=` operator, in order to reduce confusion.

- In certain contexts, the operators = and := cannot be used interchangeably; it is important to pick the correct assignment operator for the following cases:
 - When assigning a default value, you should use the = operator, as indicated in the documentation at <http://www.postgresql.org/docs/current/interactive/sql-createfunction.html>
 - For named notations in a function call, you should use the := operator

- The following example shows a case in which we cannot use = and := interchangeably:

```
car_portal=# CREATE OR REPLACE FUNCTION cast_numeric_to_int (numeric_value numeric, round boolean = TRUE /*correct use of "= ". Using ":= " will raise a syntax error */)
car_portal=# RETURNS INT AS
car_portal=# $$
car_portal$# BEGIN
car_portal$#     RETURN (CASE WHEN round = TRUE THEN CAST (numeric_value AS INTEGER)
car_portal$#                     WHEN numeric_value >= 0 THEN CAST (numeric_value -.5 AS INTEGER)
car_portal$#                     WHEN numeric_value < 0 THEN CAST (numeric_value +.5 AS INTEGER)
car_portal$#                     ELSE NULL
car_portal$#             END);
car_portal$# END;
car_portal$# $$ LANGUAGE plpgsql;
CREATE FUNCTION
car_portal=#
```

- To show how to call the `cast_numeric_to_int` function, let's execute the following code:

```
car_portal=# SELECT cast_numeric_to_int(2.3, round := true);
 cast_numeric_to_int
-----
                2
(1 row)

car_portal=# SELECT cast_numeric_to_int(2.3, round = true);
ERROR:  column "round" does not exist
LINE 1: SELECT cast_numeric_to_int(2.3, round = true);
                                         ^
```

- The assignment expression can be a single atomic value, such as `pi = 3.14`, or it can be a row, as shown in the following code snippet:

```
car_portal=# DO $$
car_portal$# DECLARE
car_portal$#     test record;
car_portal$# BEGIN
car_portal$#     test = ROW (1,'hello', 3.14);
car_portal$#     RAISE notice '%', test;
car_portal$# END;
car_portal$# $$ LANGUAGE plpgsql;
NOTICE:  (1,hello,3.14)
DO
car_portal=#
car_portal=# DO $$
car_portal$# DECLARE
car_portal$#     number_of_accounts INT :=0;
car_portal$# BEGIN
car_portal$#     number_of_accounts := (SELECT COUNT(*) FROM car_portal_app.account)::INT;
car_portal$#     RAISE NOTICE 'number_of accounts: %', number_of_accounts;
car_portal$# END;
car_portal$# $$ LANGUAGE plpgsql;
NOTICE:  number_of accounts: 481
DO
car_portal=#
```

- There are other techniques for assigning values to variables from a query that returns a single row:
 - `SELECT select_expressions INTO [STRICT] targets FROM ...;`
 - `INSERT ... RETURNING expressions INTO [STRICT] targets;`
 - `UPDATE ... RETURNING expressions INTO [STRICT] targets;`
 - `DELETE ... RETURNING expressions INTO [STRICT] targets;`

- Often, expressions are column names, while targets are variable names.
- In the case of `SELECT INTO`, the target can be of the `record` type.
- The `INSERT ... RETURNING` query is often used to return the default value of a certain column; this can be used to define the ID of a primary key, using the `SERIAL` and `BIGSERIAL` data types.

```
car_portal=# CREATE TABLE test (  
car_portal(#   id SERIAL PRIMARY KEY,  
car_portal(#   name TEXT NOT NULL  
car_portal(# );  
CREATE TABLE  
car_portal=# DO $$  
car_portal$#   DECLARE  
car_portal$#       auto_generated_id INT;  
car_portal$#   BEGIN  
car_portal$#       INSERT INTO test(name) VALUES ('Hello World') RETURNING id INTO auto_generated_id;  
car_portal$#       RAISE NOTICE 'The primary key is: %', auto_generated_id;  
car_portal$#   END  
car_portal$# $$;  
NOTICE:  The primary key is: 1  
DO  
car_portal=#
```

- You can get the default value when inserting a row in plain SQL by using CTE, as follows:

```
WITH get_id AS (  
    INSERT INTO test(name) VALUES ('Hello World')  
    RETURNING id  
)  
SELECT * FROM get_id;
```


- Finally, you can use qualified names to perform assignments; in trigger functions, you can use `NEW` and `OLD` to manipulate the values of these records.

Conditional statements

- PostgreSQL supports the `IF` and `CASE` statements, which allow for execution based on a certain condition.
- PostgreSQL supports the `IF` statement construct, as follows:
 - `IF ... THEN ... END IF`
 - `IF ... THEN ... ELSE ... END IF`
 - `IF ... THEN ... ELSEIF ... THEN ... ELSE ... END IF`
- The `CASE` statement comes in two forms, as follows:
 - `CASE ... WHEN ... THEN ... ELSE ... END CASE`
 - `CASE WHEN ... THEN ... ELSE ... END CASE`

- To understand the IF statement, let's suppose that we would like to convert the advertisement rank into text, as follows:

```
car_portal=# CREATE OR REPLACE FUNCTION cast_rank_to_text (rank int) RETURNS TEXT AS $$
car_portal$# DECLARE
car_portal$#     rank ALIAS FOR $1;
car_portal$#     rank_result TEXT;
car_portal$# BEGIN
car_portal$#     IF rank = 5 THEN rank_result = 'Excellent';
car_portal$#     ELSIF rank = 4 THEN rank_result = 'Very Good';
car_portal$#     ELSIF rank = 3 THEN rank_result = 'Good';
car_portal$#     ELSIF rank = 2 THEN rank_result = 'Fair';
car_portal$#     ELSIF rank = 1 THEN rank_result = 'Poor';
car_portal$#     ELSE rank_result = 'No such rank';
car_portal$#     END IF;
car_portal$#     RETURN rank_result;
car_portal$# END;
car_portal$# $$ Language plpgsql;
CREATE FUNCTION
car_portal=#
```

- To test the function, run the following command:

```
car_portal=# SELECT n,cast_rank_to_text(n) FROM generate_series(1,6) as foo(n);
 n | cast_rank_to_text
---+-----
 1 | Poor
 2 | Fair
 3 | Good
 4 | Very Good
 5 | Excellent
 6 | No such rank
(6 rows)
```

- The following code snippet implements the preceding function, using the CASE statement:

```
car_portal=# CREATE OR REPLACE FUNCTION cast_rank_to_text (rank int) RETURNS TEXT AS $$
car_portal$# DECLARE
car_portal$#   rank ALIAS FOR $1;
car_portal$#   rank_result TEXT;
car_portal$# BEGIN
car_portal$#   CASE rank
car_portal$#     WHEN 5 THEN rank_result = 'Excellent';
car_portal$#     WHEN 4 THEN rank_result = 'Very Good';
car_portal$#     WHEN 3 THEN rank_result = 'Good';
car_portal$#     WHEN 2 THEN rank_result = 'Fair';
car_portal$#     WHEN 1 THEN rank_result = 'Poor';
car_portal$#     ELSE rank_result = 'No such rank';
car_portal$#   END CASE;
car_portal$#   RETURN rank_result;
car_portal$# END;$$ Language plpgsql;
CREATE FUNCTION
car_portal=# SELECT n,cast_rank_to_text(n) FROM generate_series(1,6) as foo(n);
 n | cast_rank_to_text
---+-----
 1 | Poor
 2 | Fair
 3 | Good
 4 | Very Good
 5 | Excellent
 6 | No such rank
(6 rows)
```

- In the previous form of CASE, we cannot use it to match NULL values, since NULL equality is NULL.
- To overcome this limitation, you should specify the matching condition explicitly, using the second form of the CASE statement, as follows:

```
car_portal=# CREATE OR REPLACE FUNCTION cast_rank_to_text (rank int) RETURNS TEXT AS $$
car_portal$# DECLARE
car_portal$#     rank ALIAS FOR $1;
car_portal$#     rank_result TEXT;
car_portal$# BEGIN
car_portal$#     CASE
car_portal$#         WHEN rank=5 THEN rank_result = 'Excellent';
car_portal$#         WHEN rank=4 THEN rank_result = 'Very Good';
car_portal$#         WHEN rank=3 THEN rank_result = 'Good';
car_portal$#         WHEN rank=2 THEN rank_result = 'Fair';
car_portal$#         WHEN rank=1 THEN rank_result = 'Poor';
car_portal$#         WHEN rank IS NULL THEN RAISE EXCEPTION 'Rank should be not NULL';
car_portal$#         ELSE rank_result = 'No such rank';
car_portal$#     END CASE;
car_portal$#     RETURN rank_result;
car_portal$# END;
car_portal$# $$ Language plpgsql;
CREATE FUNCTION
car_portal=#
```

- To test the function, use the following command:

```
car_portal=# SELECT cast_rank_to_text(null);  
ERROR:  Rank should be not NULL  
CONTEXT:  PL/pgSQL function cast_rank_to_text(integer) line 12 at RAISE  
car_portal=#
```

- Finally, the `CASE` statement raises an exception if no branch is matched and the `ELSE` branch is not specified, as follows:

```
car_portal=# DO $$
car_portal$# DECLARE
car_portal$#   i int := 0;
car_portal$# BEGIN
car_portal$#   case WHEN i=1 then
car_portal$#       RAISE NOTICE 'i is one';
car_portal$#   END CASE;
car_portal$# END;
car_portal$# $$ LANGUAGE plpgsql;
ERROR:  case not found
HINT:   CASE statement is missing ELSE part.
CONTEXT: PL/pgSQL function inline_code_block line 5 at CASE
car_portal=#
```


Iterations

Loop statements

- The basic LOOP statement has the following structure:

```
[ <<label>> ]  
LOOP  
    statements  
END LOOP [ label ] ;
```

- To understand the `LOOP` statement, let's rewrite the `factorial` function, as follows:

```
car_portal=# DROP FUNCTION IF EXISTS factorial (int);
DROP FUNCTION
car_portal=# CREATE OR REPLACE FUNCTION factorial (fact int) RETURNS BIGINT AS $$
car_portal$# DECLARE
car_portal$#     result bigint = 1;
car_portal$# BEGIN
car_portal$#     IF fact = 1 THEN RETURN 1;
car_portal$#     ELIF fact IS NULL or fact < 1 THEN RAISE EXCEPTION 'Provide a positive integer';
car_portal$#     ELSE
car_portal$#         LOOP
car_portal$#             result = result*fact;
car_portal$#             fact = fact-1;
car_portal$#             EXIT WHEN fact = 1;
car_portal$#         END Loop;
car_portal$#     END IF;
car_portal$#     RETURN result;
car_portal$# END; $$ LANGUAGE plpgsql;
CREATE FUNCTION
car_portal=#
```

```
car_portal=# select factorial(5);
 factorial
-----
          120
(1 row)
```

The WHILE loop statement

- The `WHILE` statement keeps executing a block of statements while a particular condition is met.

- Its syntax is as follows:

```
[ <<label>> ]
```

```
WHILE boolean-expression LOOP  
    statements
```

```
END LOOP [ label ];
```

- The following example uses the `while` loop to print the days of the current month:

```
car_portal=# DO $$
car_portal$# DECLARE
car_portal$#   first_day_in_month date := date_trunc('month', current_date)::date;
car_portal$#   last_day_in_month date := (date_trunc('month', current_date)+ INTERVAL '1 MONTH - 1 day')::date;
car_portal$#   counter date = first_day_in_month;
car_portal$# BEGIN
car_portal$#   WHILE (counter <= last_day_in_month) LOOP
car_portal$#       RAISE notice '%', counter;
car_portal$#       counter := counter + interval '1 day';
car_portal$#   END LOOP;
car_portal$# END;
car_portal$# $$ LANGUAGE plpgsql;
NOTICE:  2019-10-01
NOTICE:  2019-10-02
NOTICE:  2019-10-03
NOTICE:  2019-10-04
NOTICE:  2019-10-05
NOTICE:  2019-10-06
NOTICE:  2019-10-07
NOTICE:  2019-10-08
NOTICE:  2019-10-09
NOTICE:  2019-10-10
NOTICE:  2019-10-11
NOTICE:  2019-10-12
NOTICE:  2019-10-13
NOTICE:  2019-10-14
NOTICE:  2019-10-15
```

The FOR loop statement

- PL/pgSQL provides two forms of the `FOR` statement, and they are used to execute the following:
 - Iterate through the rows returned by an SQL query
 - Iterate through a range of integer values

- **The syntax of the FOR loop statement is as follows:**

```
[ <<label>> ]  
FOR name IN [ REVERSE ] expression1 .. expression2 [ BY expression ]  
LOOP  
    statementsEND  
LOOP [ label ];
```


- The following example shows a `FOR` loop iterating over a negative range of numbers in a reverse order.

```
car_portal=# DO $$
car_portal$# BEGIN
car_portal$#   FOR j IN REVERSE -1 .. -10 BY 2 LOOP
car_portal$#     Raise notice '%', j;
car_portal$#   END LOOP;
car_portal$# END; $$ LANGUAGE plpgsql;
NOTICE:  -1
NOTICE:  -3
NOTICE:  -5
NOTICE:  -7
NOTICE:  -9
DO
car_portal=#
```

- To iterate through the results of a set query, the syntax is different, as follows:

```
[ <<label>> ]  
FOR target IN query LOOP  
    statements  
END LOOP [ label ] ;
```

- The following example shows all of the database names:

```
car_portal=# DO $$
car_portal$# DECLARE
car_portal$#   database RECORD;
car_portal$# BEGIN
car_portal$#   FOR database IN SELECT * FROM pg_database LOOP
car_portal$#     RAISE notice '%', database.datname;
car_portal$#   END LOOP;
car_portal$# END; $$;
NOTICE: postgres
NOTICE: car_portal
NOTICE: template1
NOTICE: template0
DO
car_portal=#
```

Returning from the function

Returning voids

- A `void` type is used in a function to perform some side effects, such as logging; the built-in function, `pg_sleep`, is used to delay the execution of a server process, in seconds, as follows:

```
car_portal=# \df pg_sleep
              List of functions
 Schema | Name | Result data type | Argument data types | Type 
-----+-----+-----+-----+-----
 pg_catalog | pg_sleep | void | double precision | func
(1 row)
```

Returning a single row

- The RETURN type can be a base, composite, table, pseudo, or domain data type.
- The following function returns a JSON representation of a certain account:

```
car_portal=# CREATE OR REPLACE FUNCTION car_portal_app.get_account_in_json (account_id INT) RETURNS JSON AS $$
car_portal$#   SELECT row_to_json(account) FROM car_portal_app.account WHERE account_id = $1;
car_portal$# $$ LANGUAGE SQL;
CREATE FUNCTION
car_portal=# CREATE OR REPLACE FUNCTION car_portal_app.get_account_in_json1 (acc_id INT) RETURNS JSON AS $$
car_portal$# BEGIN
car_portal$#   RETURN (SELECT row_to_json(account) FROM car_portal_app.account WHERE account_id = acc_id);
car_portal$# END;
car_portal$# $$ LANGUAGE plpgsql;
CREATE FUNCTION
```

```
car_portal=# select car_portal_app.get_account_in_json(1);
```

```
get_account_in_json
```

```
-----  
-----  
{ "account_id":1,"first_name":"James","last_name":"Butt","email":"jbutt@gmail.com","password":"1b9ef408e82e38346e6ebbf2dcc5ece"}  
(1 row)
```

```
car_portal=# select car_portal_app.get_account_in_json1(1);
```

```
get_account_in_json1
```

```
-----  
-----  
{ "account_id":1,"first_name":"James","last_name":"Butt","email":"jbutt@gmail.com","password":"1b9ef408e82e38346e6ebbf2dcc5ece"}  
(1 row)
```

Returning multiple rows

- **Set returning functions (SRFs)** can be used to return a set of rows.
- The row type can either be a base type, such as an integer, composite, table type, pseudo type, or domain type.
- To return a set from a function, the `SETOF` keyword is used to mark the function as an SRF, as follows:

```
car_portal=# CREATE OR REPLACE FUNCTION car_portal_app.car_model(model_name TEXT) RETURNS SETOF car_portal_app.car_model AS
$$
car_portal$#   SELECT car_model_id, make, model FROM car_portal_app.car_model WHERE model = model_name;
car_portal$# $$ LANGUAGE SQL;
CREATE FUNCTION
car_portal=# CREATE OR REPLACE FUNCTION car_portal_app.car_model1(model_name TEXT) RETURNS SETOF car_portal_app.car_model A
S $$
car_portal$# BEGIN
car_portal$#   RETURN QUERY SELECT car_model_id, make, model FROM car_portal_app.car_model WHERE model = model_name;
car_portal$# END;
car_portal$# $$ LANGUAGE plpgsql;
CREATE FUNCTION
```


- To test the previous functions, let's run them.
- Note the caching effect of `plpgsql` on the performance:

```
car_portal=# \timing
Timing is on.
car_portal=# SELECT * FROM car_portal_app.car_model('A1');
 car_model_id | make | model 
-----+-----+-----
           1 | Audi | A1
(1 row)

Time: 0.911 ms
car_portal=# SELECT * FROM car_portal_app.car_model1('A1');
 car_model_id | make | model 
-----+-----+-----
           1 | Audi | A1
(1 row)

Time: 0.807 ms
```

- If the return type is not defined, you can do the following:
 - Define a new data type and use it
 - Use a return table
 - Use output parameters and record the data type

- Let's suppose that we would only like to return `car_model_id` and `make`; as in this case, we do not have a data type defined.
- The preceding function could be written as follows:

```
car_portal=# CREATE OR REPLACE FUNCTION car_portal_app.car_model2(model_name TEXT) RETURNS TABLE (car_model_id INT , make TEXT) AS $$
car_portal$# SELECT car_model_id, make FROM car_portal_app.car_model WHERE model = model_name;
car_portal$# $$ LANGUAGE SQL;
CREATE FUNCTION
car_portal=# CREATE OR REPLACE FUNCTION car_portal_app.car_model3(model_name TEXT) RETURNS TABLE (car_model_id INT , make TEXT) AS $$
car_portal$# BEGIN
car_portal$# RETURN QUERY SELECT car_model_id, make FROM car_portal_app.car_model WHERE model = model_name;
car_portal$# END;
car_portal$# $$ LANGUAGE plpgsql;
CREATE FUNCTION
```

- To test the `car_model2` and `car_model3` functions, let's run the following code:

```
car_portal=# SELECT * FROM car_portal_app.car_model2('A1');
 car_model_id | make
-----+-----
           1 | Audi
(1 row)

car_portal=# SELECT * FROM car_portal_app.car_model3('A1');
ERROR:  column reference "car_model_id" is ambiguous
LINE 1: SELECT car_model_id, make FROM car_portal_app.car_model WHERE...
              ^
DETAIL:  It could refer to either a PL/pgSQL variable or a table column.
QUERY:  SELECT car_model_id, make FROM car_portal_app.car_model WHERE model = model_name
CONTEXT:  PL/pgSQL function car_portal_app.car_model3(text) line 3 at RETURN QUERY
car_portal=#
```

- Note that, in the preceding function, an error is raised, because `plpgsql` was confusing the column name with the table definition.
- The reason is that the return table is shorthand for writing `OUTPUT` parameters.
- To fix this, we need to rename the attribute names, as follows:

```
car_portal=# CREATE OR REPLACE FUNCTION car_portal_app.car_model3(model_name TEXT) RETURNS TABLE (car_model_id INT , make TEXT) AS $$
car_portal$# BEGIN
car_portal$#   RETURN QUERY SELECT a.car_model_id, a.make FROM car_portal_app.car_model a WHERE model = model_name;
car_portal$# END;
car_portal$# $$ LANGUAGE plpgsql;
CREATE FUNCTION
car_portal=# SELECT * FROM car_portal_app.car_model3('A1');
 car_model_id | make
-----+-----
            1 | Audi
(1 row)
```

- The preceding function can also be written using the OUT variables, as follows:

```
car_portal=# CREATE OR REPLACE FUNCTION car_portal_app.car_model4(model_name TEXT, OUT car_model_id INT, OUT make TEXT ) RE
TURNS SETOF RECORD AS $$
car_portal$# BEGIN
car_portal$#   RETURN QUERY SELECT a.car_model_id, a.make FROM car_portal_app.car_model a WHERE model = model_name;
car_portal$# END;
car_portal$# $$ LANGUAGE plpgsql;
CREATE FUNCTION
car_portal=# SELECT * FROM car_portal_app.car_model4('A1'::text);
 car_model_id | make
-----+-----
            1 | Audi
(1 row)
```