

The Relational Data Model and Relational Database Constraints

Part 2

Relational Model Constraints and Relational Database Schemas

Part 1

- Constraints on databases can generally be divided into three main categories:
 - Constraints that are inherent in the data model. We call these **inherent model-based constraints** or **implicit constraints**.
 - Constraints that can be directly expressed in schemas of the data model, typically by specifying them in the DDL (data definition language). We call these **schema-based constraints** or **explicit constraints**.
 - Constraints that *cannot* be directly expressed in the schemas of the data model, and hence must be expressed and enforced by the application programs. We call these **application-based** or **semantic constraints** or business rules.

- The characteristics of relations that we discussed earlier are the inherent constraints of the relational model and belong to the first category.
- For example, the constraint that a relation cannot have duplicate tuples is an inherent constraint.

- The schema-based constraints include domain constraints, key constraints, constraints on NULLs, entity integrity constraints, and referential integrity constraints.

Domain Constraints

- Domain constraints specify that within each tuple, the value of each attribute A must be an atomic value from the domain $\text{dom}(A)$.

Key Constraints and Constraints on NULL Values

- A **superkey** SK specifies a *uniqueness constraint* that no two distinct tuples in any state r of R can have the same value for SK.
- Every relation has at least one default superkey—the set of all its attributes.

- A superkey can have redundant attributes, however, so a more useful concept is that of a *key*, which has no redundancy.

- A **key** K of a relation schema R is a superkey of R with the additional property that removing any attribute A from K leaves a set of attributes K' that is not a superkey of R any more.
- Hence, a key satisfies two properties:
 1. Two distinct tuples in any state of the relation cannot have identical values for (all) the attributes in the key. This first property also applies to a superkey.
 2. It is a *minimal superkey*—that is, a superkey from which we cannot remove any attributes and still have the uniqueness constraint in condition 1 hold. This property is not required by a superkey.

- Notice that a set of attributes constituting a key is a property of the relation schema; it is a constraint that should hold on *every* valid relation state of the schema.
- A key is determined from the meaning of the attributes, and the property is *time-invariant*: It must continue to hold when we insert new tuples in the relation.

- In general, a relation schema may have more than one key.
- In this case, each of the keys is called a **candidate key**.
- It is common to designate one of the candidate keys as the **primary key** of the relation.
- We use the convention that the attributes that form the primary key of a relation schema are underlined.

CAR

| <u>License_number</u> | Engine_serial_number | Make | Model | Year |
|-----------------------|----------------------|------------|---------|------|
| Texas ABC-739 | A69352 | Ford | Mustang | 02 |
| Florida TVP-347 | B43696 | Oldsmobile | Cutlass | 05 |
| New York MPO-22 | X83554 | Oldsmobile | Delta | 01 |
| California 432-TFY | C43742 | Mercedes | 190-D | 99 |
| California RSK-629 | Y82935 | Toyota | Camry | 04 |
| Texas RSK-629 | U028365 | Jaguar | XJS | 04 |

- Notice that when a relation schema has several candidate keys, the choice of one to become the primary key is somewhat arbitrary; however, it is usually better to choose a primary key with a single attribute or a small number of attributes.
- The other candidate keys are designated as **unique keys**, and are not underlined.

- Another constraint on attributes specifies whether NULL values are or are not permitted.