

Basic SQL

Part 1

SQL Data Definition and Data Types

- SQL uses the terms **table**, **row**, and **column** for the formal relational model terms *relation*, *tuple*, and *attribute*, respectively.

Schema and Catalog Concepts in SQL

- An **SQL schema** is identified by a **schema name**, and includes an **authorization identifier** to indicate the user or account who owns the schema, as well as **descriptors** for *each element* in the schema.
- Schema **elements** include tables, constraints, views, domains, and other constructs (such as authorization grants) that describe the schema.

- A schema is created via the `CREATE SCHEMA` statement, which can include all the schema elements' definitions.
- Alternatively, the schema can be assigned a name and authorization identifier, and the elements can be defined later.
- For example, the following statement creates a schema called `COMPANY`, owned by the user with authorization identifier `'Jsmith'`.
 - Note that each statement in SQL ends with a semicolon.
 - **`CREATE SCHEMA COMPANY AUTHORIZATION 'Jsmith';`**

- In general, not all users are authorized to create schemas and schema elements.
- The privilege to create schemas, tables, and other constructs must be explicitly granted to the relevant user accounts by the system administrator or DBA.

- In addition to the concept of a schema, SQL uses the concept of a **catalog**—a named collection of schemas in an SQL environment.
- An SQL **environment** is basically an installation of an SQL-compliant RDBMS on a computer system.
- A catalog always contains a special schema called `INFORMATION_SCHEMA`, which provides information on all the schemas in the catalog and all the element descriptors in these schemas.
- Integrity constraints such as referential integrity can be defined between relations only if they exist in schemas within the same catalog.
- Schemas within the same catalog can also share certain elements, such as domain definitions.

The CREATE TABLE Command in SQL


```

CREATE TABLE EMPLOYEE
( Fname          VARCHAR(15)          NOT NULL,
  Minit          CHAR,
  Lname          VARCHAR(15)          NOT NULL,
  Ssn            CHAR(9)              NOT NULL,
  Bdate          DATE,
  Address        VARCHAR(30),
  Sex            CHAR,
  Salary         DECIMAL(10,2),
  Super_ssn      CHAR(9),
  Dno            INT                  NOT NULL,
  PRIMARY KEY (Ssn),
  FOREIGN KEY (Super_ssn) REFERENCES EMPLOYEE(Ssn),
  FOREIGN KEY (Dno) REFERENCES DEPARTMENT(Dnumber) );

CREATE TABLE DEPARTMENT
( Dname          VARCHAR(15)          NOT NULL,
  Dnumber        INT                  NOT NULL,
  Mgr_ssn        CHAR(9)              NOT NULL,
  Mgr_start_date DATE,
  PRIMARY KEY (Dnumber),
  UNIQUE (Dname),
  FOREIGN KEY (Mgr_ssn) REFERENCES EMPLOYEE(Ssn) );

CREATE TABLE DEPT_LOCATIONS
( Dnumber        INT                  NOT NULL,
  Dlocation      VARCHAR(15)          NOT NULL,
  PRIMARY KEY (Dnumber, Dlocation),
  FOREIGN KEY (Dnumber) REFERENCES DEPARTMENT(Dnumber) );

CREATE TABLE PROJECT
( Pname          VARCHAR(15)          NOT NULL,
  Pnumber        INT                  NOT NULL,
  Plocation      VARCHAR(15),
  Dnum           INT                  NOT NULL,
  PRIMARY KEY (Pnumber),
  UNIQUE (Pname),
  FOREIGN KEY (Dnum) REFERENCES DEPARTMENT(Dnumber) );

CREATE TABLE WORKS_ON
( Essn           CHAR(9)              NOT NULL,
  Pno            INT                  NOT NULL,
  Hours          DECIMAL(3,1)         NOT NULL,
  PRIMARY KEY (Essn, Pno),
  FOREIGN KEY (Essn) REFERENCES EMPLOYEE(Ssn),
  FOREIGN KEY (Pno) REFERENCES PROJECT(Pnumber) );

CREATE TABLE DEPENDENT
( Essn           CHAR(9)              NOT NULL,
  Dependent_name VARCHAR(15)          NOT NULL,
  Sex            CHAR,
  Bdate          DATE,
  Relationship    VARCHAR(8),
  PRIMARY KEY (Essn, Dependent_name),
  FOREIGN KEY (Essn) REFERENCES EMPLOYEE(Ssn) );

```

- Typically, the SQL schema in which the relations are declared is implicitly specified in the environment in which the CREATE TABLE statements are executed.
- Alternatively, we can explicitly attach the schema name to the relation name, separated by a period.
 - **CREATE TABLE** COMPANY.EMPLOYEE . . .

- The relations declared through `CREATE TABLE` statements are called **base tables** (or base relations); this means that the relation and its tuples are actually created and stored as a file by the DBMS.
- Base relations are distinguished from **virtual relations**, created through the **CREATE VIEW** statement, which may or may not correspond to an actual physical file.
- In SQL, the attributes in a base table are considered to be *ordered in the sequence in which they are specified* in the `CREATE TABLE` statement.
- However, rows (tuples) are not considered to be ordered within a relation.

Attribute Data Types and Domains in SQL

- The basic **data types** available for attributes include numeric, character string, bit string, Boolean, date, and time.

- It is possible to specify the data type of each attribute directly; alternatively, a domain can be declared, and the domain name used with the attribute specification.
- This makes it easier to change the data type for a domain that is used by numerous attributes in a schema, and improves schema readability.
- For example, we can create a domain SSN_TYPE by the following statement:

```
CREATE DOMAIN SSN_TYPE AS CHAR(9) ;
```