

More SQL: Complex Queries, Triggers, Views, and Schema Modification

Part 1

More Complex SQL Retrieval Queries

Comparisons Involving NULL and Three-Valued Logic

- In general, each individual `NULL` value is considered to be different from every other `NULL` value in the various database records.
- When a `NULL` is involved in a comparison operation, the result is considered to be `UNKNOWN` (it may be `TRUE` or it may be `FALSE`).
- Hence, SQL uses a three-valued logic with values `TRUE`, `FALSE`, and `UNKNOWN` instead of the standard two-valued (Boolean) logic with values `TRUE` or `FALSE`.
- It is therefore necessary to define the results (or truth values) of three-valued logical expressions when the logical connectives `AND`, `OR`, and `NOT` are used.

(a)	AND	TRUE	FALSE	UNKNOWN
	TRUE	TRUE	FALSE	UNKNOWN
	FALSE	FALSE	FALSE	FALSE
	UNKNOWN	UNKNOWN	FALSE	UNKNOWN
(b)	OR	TRUE	FALSE	UNKNOWN
	TRUE	TRUE	TRUE	TRUE
	FALSE	TRUE	FALSE	UNKNOWN
	UNKNOWN	TRUE	UNKNOWN	UNKNOWN
(c)	NOT			
	TRUE	FALSE		
	FALSE	TRUE		
	UNKNOWN	UNKNOWN		

- In select-project-join queries, the general rule is that only those combinations of tuples that evaluate the logical expression in the WHERE clause of the query to TRUE are selected.
- Tuple combinations that evaluate to FALSE or UNKNOWN are not selected.
- However, there are exceptions to that rule for certain operations, such as outer joins.

- SQL allows queries that check whether an attribute value is `NULL`.
- Rather than using `=` or `<>` to compare an attribute value to `NULL`, SQL uses the comparison operators `IS` or `IS NOT`.
- This is because SQL considers each `NULL` value as being distinct from every other `NULL` value, so equality comparison is not appropriate.
- It follows that when a join condition is specified, tuples with `NULL` values for the join attributes are not included in the result (unless it is an `OUTER JOIN`).

- **Query 18.** Retrieve the names of all employees who do not have supervisors.
- Q18: SELECT Fname, Lname
 FROM EMPLOYEE
 WHERE Super_ssn IS NULL;

Nested Queries, Tuples, and Set/Multiset Comparisons

- **Query 4.** Make a list of all project numbers for projects that involve an employee whose last name is 'Smith', either as a worker or as a manager of the department that controls the project.

- **Q4A:** SELECT DISTINCT Pnumber
FROM PROJECT
WHERE Pnumber IN (SELECT Pnumber
FROM PROJECT, DEPARTMENT,
EMPLOYEE
WHERE Dnum=Dnumber AND
Mgr_ssn=Ssn AND
Lname='Smith')

OR
Pnumber IN (SELECT Pno
FROM WORKS_ON, EMPLOYEE
WHERE Essn=Ssn AND
Lname='Smith');

- SQL allows the use of **tuples** of values in comparisons by placing them within parentheses.
- ```
SELECT DISTINCT Essn
FROM WORKS_ON
WHERE (Pno, Hours) IN (SELECT Pno, Hours
 FROM WORKS_ON
 WHERE Essn='123456789');
```

- ```
SELECT Lname, Fname
FROM EMPLOYEE
WHERE Salary > ALL (SELECT Salary
                     FROM EMPLOYEE
                     WHERE Dno=5) ;
```

- **Query 16.** Retrieve the name of each employee who has a dependent with the same first name and is the same sex as the employee.
- **Q16:**

```
SELECT E.Fname, E.Lname
FROM   EMPLOYEE AS E
WHERE  E.Ssn IN (SELECT Essn
                FROM   DEPENDENT AS D
                WHERE  E.Fname=D.Dependent_name
                AND
                E.Sex=D.Sex) ;
```

- It is generally advisable to create tuple variables (aliases) for *all the tables referenced in an SQL query* to avoid potential errors and ambiguities.

Correlated Nested Queries

- Whenever a condition in the `WHERE` clause of a nested query references some attribute of a relation declared in the outer query, the two queries are said to be **correlated**.
- We can understand a correlated query better by considering that the *nested query is evaluated once for each tuple (or combination of tuples) in the outer query.*

- In general, a query written with nested select-from-where blocks and using the = or IN comparison operators can *always* be expressed as a single block query.
- Q16A:

```
SELECT E.Fname, E.Lname
FROM   EMPLOYEE AS E, DEPENDENT AS D
WHERE  E.Ssn=D.Essn AND
       E.Sex=D.Sex AND
       E.Fname=D.Dependent_name;
```

The EXISTS and UNIQUE Functions in SQL

- The `EXISTS` function in SQL is used to check whether the result of a correlated nested query is *empty* (contains no tuples) or not.
- The result of `EXISTS` is a Boolean value **TRUE** if the nested query result contains at least one tuple, or **FALSE** if the nested query result contains no tuples.

- **Query 6.** Retrieve the names of employees who have no dependents.

- Q6:

```
SELECT Fname, Lname
FROM EMPLOYEE
WHERE NOT EXISTS (SELECT *
                   FROM DEPENDENT
                   WHERE Ssn=Essn) ;
```

- **Query 7.** List the names of managers who have at least one dependent.

- **Q7:**

```
SELECT  Fname, Lname
FROM    EMPLOYEE
WHERE   EXISTS (SELECT *
                  FROM    DEPENDENT
                  WHERE    Ssn=Essn)

AND
EXISTS  (SELECT *
          FROM    DEPARTMENT
          WHERE    Ssn=Mgr_ssn ) ;
```

- **Q3:** Retrieve the name of each employee who works on all the projects controlled by department number 5.
- **Q3A:**

```
SELECT Fname, Lname
FROM   EMPLOYEE
WHERE  NOT EXISTS ( (SELECT Pnumber
                      FROM   PROJECT
                      WHERE   Dnum=5)
                  EXCEPT
                  (SELECT Pno
                      FROM   WORKS_ON
                      WHERE   Ssn=Essn) ) ;
```

- Q3B: SELECT Lname, Fname
FROM EMPLOYEE
WHERE NOT EXISTS (SELECT *
FROM WORKS_ON B
WHERE (B.Pno IN (SELECT Pnumber
FROM PROJECT
WHERE Dnum=5)

AND
NOT EXISTS (SELECT *
FROM WORKS_ON C
WHERE C.Essn=Ssn AND
C.Pno=B.Pno))) ;

- There is another SQL function, `UNIQUE (Q)` , which returns `TRUE` if there are no duplicate tuples in the result of query `Q`; otherwise, it returns `FALSE`.
- This can be used to test whether the result of a nested query is a set or a multiset.

Explicit Sets and Renaming of Attributes in SQL

- **Query 17.** Retrieve the Social Security numbers of all employees who work on project numbers 1, 2, or 3.
- Q17:

```
SELECT DISTINCT Essn
FROM   WORKS_ON
WHERE  Pno IN (1, 2, 3);
```

- **Q8A:** SELECT E.Lname AS Employee_name,
 S.Lname AS Supervisor_name
FROM EMPLOYEE AS E, EMPLOYEE AS S
WHERE E.Super_ssn=S.Ssn;

Joined Tables in SQL and Outer Joins

- Q1A:

```
SELECT Fname, Lname, Address
FROM    (EMPLOYEE JOIN DEPARTMENT
          ON Dno=Dnumber)
WHERE   Dname='Research';
```

- **Q1B:** SELECT Fname, Lname, Address
FROM (EMPLOYEE NATURAL JOIN
(DEPARTMENT AS DEPT (Dname, Dno, Mssn, Msdate)))
WHERE Dname='Research';

- **Q8B:** SELECT E.Lname AS Employee_name,
 S.Lname AS Supervisor_name
FROM (EMPLOYEE AS E LEFT OUTER JOIN
 EMPLOYEE AS S
 ON E.Super_ssn=S.Ssn) ;

- There are a variety of outer join operations.
- In SQL, the options available for specifying joined tables include `INNER JOIN` (only pairs of tuples that match the join condition are retrieved, same as `JOIN`), `LEFT OUTER JOIN` (every tuple in the left table must appear in the result; if it does not have a matching tuple, it is padded with `NULL` values for the attributes of the right table), `RIGHT OUTER JOIN` (every tuple in the right table must appear in the result; if it does not have a matching tuple, it is padded with `NULL` values for the attributes of the left table), and `FULL OUTER JOIN`.
- In the latter three options, the keyword `OUTER` may be omitted.
- If the join attributes have the same name, one can also specify the natural join variation of outer joins by using the keyword `NATURAL` before the operation (for example, `NATURAL LEFT OUTER JOIN`).
- The keyword `CROSS JOIN` is used to specify the `CARTESIAN PRODUCT` operation, although this should be used only with the utmost care because it generates all possible tuple combinations.

- Q2A:

```
SELECT Pnumber, Dnum, Lname, Address, Bdate
FROM    ((PROJECT JOIN DEPARTMENT
          ON Dnum=Dnumber)
         JOIN EMPLOYEE
          ON Mgr_ssn=Ssn)
WHERE   Plocation='Stafford';
```

Aggregate Functions in SQL

- **Aggregate functions** are used to summarize information from multiple tuples into a single-tuple summary.
- **Grouping** is used to create subgroups of tuples before summarization.
- A number of built-in aggregate functions exist: **COUNT**, **SUM**, **MAX**, **MIN**, and **AVG**.
- These functions can be used in the `SELECT` clause or in a `HAVING` clause.

- **Query 19.** Find the sum of the salaries of all employees, the maximum salary, the minimum salary, and the average salary.
- Q19: SELECT SUM (Salary), MAX (Salary),
 MIN (Salary), AVG (Salary)
 FROM EMPLOYEE;

- **Query 20.** Find the sum of the salaries of all employees of the 'Research' department, as well as the maximum salary, the minimum salary, and the average salary in this department.

- **Q20:**

```
SELECT SUM (Salary), MAX (Salary),  
        MIN (Salary), AVG (Salary)  
FROM    (EMPLOYEE JOIN DEPARTMENT  
          ON Dno=Dnumber)  
WHERE   Dname='Research';
```


- **Queries 21 and 22.** Retrieve the total number of employees in the company (Q21) and the number of employees in the 'Research' department (Q22).
- Q21: SELECT COUNT (*)
 FROM EMPLOYEE;
- Q22: SELECT COUNT (*)
 FROM EMPLOYEE, DEPARTMENT
 WHERE DNO=DNUMBER AND
 DNAME='Research' ;

- **Query 23.** Count the number of distinct salary values in the database.
- Q23: SELECT COUNT (DISTINCT Salary)
 FROM EMPLOYEE;

- **Query 5:** Retrieve the names of all employees who have two or more dependents.
- Q5:

```
SELECT Lname, Fname
FROM EMPLOYEE
WHERE (SELECT COUNT (*)
      FROM DEPENDENT
      WHERE Ssn=Essn) >= 2;
```

Grouping: The GROUP BY and HAVING Clauses

- **Query 24.** For each department, retrieve the department number, the number of employees in the department, and their average salary.
- Q24:

```
SELECT      Dno, COUNT (*), AVG (Salary)
FROM        EMPLOYEE
GROUP BY    Dno;
```

- **Query 25.** For each project, retrieve the project number, the project name, and the number of employees who work on that project.
- Q25:

```
SELECT      Pnumber, Pname, COUNT (*)  
FROM        PROJECT, WORKS_ON  
WHERE       Pnumber=Pno  
GROUP BY    Pnumber, Pname;
```

- **Query 26.** For each project on which more than two employees work, retrieve the project number, the project name, and the number of employees who work on the project.

- Q26: SELECT Pnumber, Pname, COUNT (*)
 FROM PROJECT, WORKS_ON
 WHERE Pnumber=Pno
 GROUP BY Pnumber, Pname
 HAVING COUNT (*) > 2;

- **Query 27.** For each project, retrieve the project number, the project name, and the number of employees from department 5 who work on the project.

- Q27: SELECT Pnumber, Pname, COUNT (*)
 FROM PROJECT, WORKS_ON, EMPLOYEE
 WHERE Pnumber=Pno AND Ssn=Essn AND
 Dno=5
 GROUP BY Pnumber, Pname;

- **Query 28.** For each department that has more than five employees, retrieve the department number and the number of its employees who are making more than \$40,000.

• Q28:

```
SELECT      Dnumber, COUNT (*)
FROM        DEPARTMENT, EMPLOYEE
WHERE       Dnumber=Dno AND
            Salary>40000 AND
            Dno in (SELECT      Dno
                     FROM        EMPLOYEE
                     GROUP BY    Dno
                     HAVING      COUNT (*) > 5)
GROUP BY    Dnumber;
```