

# The Relational Algebra

Part 2

# Relational Algebra Operations from Set Theory

# The UNION, INTERSECTION, and MINUS Operations

- **UNION, INTERSECTION, and SET DIFFERENCE** (also called **MINUS** or **EXCEPT**) are binary operations; that is, each is applied to two sets (of tuples).
- When these operations are adapted to relational databases, the two relations on which any of these three operations are applied must have the same **type of tuples**; this condition has been called *union compatibility* or *type compatibility*.
- Two relations  $R(A_1, A_2, \dots, A_n)$  and  $S(B_1, B_2, \dots, B_n)$  are said to be **union compatible** (or **type compatible**) if they have the same degree  $n$  and if  $\text{dom}(A_i) = \text{dom}(B_i)$  for  $1 \leq i \leq n$ .

The set operations UNION, INTERSECTION, and MINUS. (a) Two union-compatible relations.  
 (b)  $\text{STUDENT} \cup \text{INSTRUCTOR}$ . (c)  $\text{STUDENT} \cap \text{INSTRUCTOR}$ . (d)  $\text{STUDENT} - \text{INSTRUCTOR}$ .  
 (e)  $\text{INSTRUCTOR} - \text{STUDENT}$ .

**(a) STUDENT**

Fn	Ln
Susan	Yao
Ramesh	Shah
Johnny	Kohler
Barbara	Jones
Amy	Ford
Jimmy	Wang
Ernest	Gilbert

**INSTRUCTOR**

Fname	Lname
John	Smith
Ricardo	Browne
Susan	Yao
Francis	Johnson
Ramesh	Shah

**(b)**

Fn	Ln
Susan	Yao
Ramesh	Shah
Johnny	Kohler
Barbara	Jones
Amy	Ford
Jimmy	Wang
Ernest	Gilbert
John	Smith
Ricardo	Browne
Francis	Johnson

**(c)**

Fn	Ln
Susan	Yao
Ramesh	Shah

**(d)**

Fn	Ln
Johnny	Kohler
Barbara	Jones
Amy	Ford
Jimmy	Wang
Ernest	Gilbert

**(e)**

Fname	Lname
John	Smith
Ricardo	Browne
Francis	Johnson

- Notice that both UNION and INTERSECTION are *commutative* operations; that is,

$$R \cup S = S \cup R \text{ and } R \cap S = S \cap R$$

- Both UNION and INTERSECTION can be treated as *n*-ary operations applicable to any number of relations because both are also *associative* operations; that is,

$$R \cup (S \cup T) = (R \cup S) \cup T \text{ and} \\ (R \cap S) \cap T = R \cap (S \cap T)$$

- The MINUS operation is not commutative; that is, in general,

$$R - S \neq S - R$$

- Note that INTERSECTION can be expressed in terms of union and set difference as follows:

$$R \cap S = ( (R \cup S) - (R - S) ) - (S - R)$$

- $\text{DEP5\_EMPS} \leftarrow \sigma_{\text{Dno}=5}(\text{EMPLOYEE})$   
 $\text{RESULT1} \leftarrow \pi_{\text{Ssn}}(\text{DEP5\_EMPS})$   
 $\text{RESULT2}(\text{Ssn}) \leftarrow \pi_{\text{Super\_ssn}}(\text{DEP5\_EMPS})$   
 $\text{RESULT} \leftarrow \text{RESULT1} \cup \text{RESULT2}$

- In SQL, there are three operations—UNION, INTERSECT, and EXCEPT—that correspond to the set operations described here.
- In addition, there are multiset operations (UNION ALL, INTERSECT ALL, and EXCEPT ALL) that do not eliminate duplicates.



# The CARTESIAN PRODUCT (CROSS PRODUCT) Operation

- The **CARTESIAN PRODUCT** operation—also known as **CROSS PRODUCT** or **CROSS JOIN**—is denoted by  $\times$ .
- In general, the result of  $R(A_1, A_2, \dots, A_n) \times S(B_1, B_2, \dots, B_m)$  is a relation  $Q$  with degree  $n + m$  attributes  $Q(A_1, A_2, \dots, A_n, B_1, B_2, \dots, B_m)$ , in that order.
- If  $R$  has  $n_R$  tuples (denoted as  $|R| = n_R$ ), and  $S$  has  $n_S$  tuples, then  $R \times S$  will have  $n_R * n_S$  tuples.

- The  $n$ -ary CARTESIAN PRODUCT operation is an extension of the above concept, which produces new tuples by concatenating all possible combinations of tuples from  $n$  underlying relations.

- In general, the CARTESIAN PRODUCT operation applied by itself is generally meaningless.
- It is mostly useful when followed by a selection that matches values of attributes coming from the component relations.
- $$\begin{aligned} \text{FEMALE\_EMPS} &\leftarrow \sigma_{\text{Sex}='F'}(\text{EMPLOYEE}) \\ \text{EMP\_NAMES} &\leftarrow \pi_{\text{Fname}, \text{Lname}, \text{Ssn}}(\text{FEMALE\_EMPS}) \\ \text{EMP\_DEPENDENTS} &\leftarrow \text{EMP\_NAMES} \times \text{DEPENDENT} \\ \text{ACTUAL\_DEPENDENTS} &\leftarrow \sigma_{\text{Ssn}=\text{Essn}}(\text{EMP\_DEPENDENTS}) \\ \text{RESULT} &\leftarrow \pi_{\text{Fname}, \text{Lname}, \text{Dependent\_name}}(\text{ACTUAL\_DEPENDENTS}) \end{aligned}$$

- In SQL, `CARTESIAN PRODUCT` can be realized by using the `CROSS JOIN` option in joined tables.
- Alternatively, if there are two tables in the `WHERE` clause and there is no corresponding join condition in the query, the result will also be the `CARTESIAN PRODUCT` of the two tables.

# Binary Relational Operations: JOIN and DIVISION

# The JOIN Operation

- The JOIN operation can be specified as a CARTESIAN PRODUCT operation followed by a SELECT operation.
- The general form of a JOIN operation on two relations  $R(A_1, A_2, \dots, A_n)$  and  $S(B_1, B_2, \dots, B_m)$  is
$$R \bowtie_{\langle \text{join condition} \rangle} S$$
- The result of the JOIN is a relation  $Q$  with  $n + m$  attributes  $Q(A_1, A_2, \dots, A_n, B_1, B_2, \dots, B_m)$  in that order;  $Q$  has one tuple for each combination of tuples—one from  $R$  and one from  $S$ —*whenever the combination satisfies the join condition.*

- A general join condition is of the form  
 $\langle \text{condition} \rangle \text{ AND } \langle \text{condition} \rangle \text{ AND } \dots \text{ AND } \langle \text{condition} \rangle$   
where each  $\langle \text{condition} \rangle$  is of the form  $A_i \theta B_j$ ,  $A_i$  is an attribute of  $R$ ,  $B_j$  is an attribute of  $S$ ,  $A_i$  and  $B_j$  have the same domain, and  $\theta$  (theta) is one of the comparison operators  $\{=, <, \leq, >, \geq, \neq\}$ .
- A JOIN operation with such a general join condition is called a **THETA JOIN**.

- $\text{DEPT\_MGR} \leftarrow \text{DEPARTMENT} \bowtie_{\text{Mgr\_ssn}=\text{Ssn}} \text{EMPLOYEE}$   
 $\text{RESULT} \leftarrow \Pi_{\text{Dname}, \text{Lname}, \text{Fname}} (\text{DEPT\_MGR})$



# Variations of JOIN: The EQUIJOIN and NATURAL JOIN

- The most common use of JOIN involves join conditions with equality comparisons only.
- Such a JOIN, where the only comparison operator used is =, is called an **EQUIJOIN**.
- Notice that in the result of an EQUIJOIN we always have one or more pairs of attributes that have identical values in every tuple.
- Because one of each pair of attributes with identical values is superfluous, a new operation called **NATURAL JOIN**—denoted by  $*$ —was created to get rid of the second (superfluous) attribute in an EQUIJOIN condition.
- The standard definition of NATURAL JOIN requires that the two join attributes (or each pair of join attributes) have the same name in both relations.
- If this is not the case, a renaming operation is applied first.

- PROJ\_DEPT  $\leftarrow$  PROJECT \*  $\rho$  (Dname, Dnum, Mgr\_ssn, Mgr\_start\_date) (DEPARTMENT)
- DEPT\_LOCS  $\leftarrow$  DEPARTMENT \* DEPT\_LOCATIONS

- In general, the join condition for NATURAL JOIN is constructed by equating *each pair of join attributes* that have the same name in the two relations and combining these conditions with **AND**.

- Notice that if no combination of tuples satisfies the join condition, the result of a JOIN is an empty relation with zero tuples.
- In general, if  $R$  has  $n_R$  tuples and  $S$  has  $n_S$  tuples, the result of a JOIN operation  $R \bowtie_{\langle \text{join condition} \rangle} S$  will have between zero and  $n_R * n_S$  tuples.
- The expected size of the join result divided by the maximum size  $n_R * n_S$  leads to a ratio called **join selectivity**, which is a property of each join condition.
- If there is no join condition, all combinations of tuples qualify and the JOIN degenerates into a CARTESIAN PRODUCT, also called CROSS PRODUCT or CROSS JOIN.

- In SQL, JOIN can be realized in several different ways.
- The first method is to specify the `<join conditions>` in the `WHERE` clause, along with any other selection conditions.
- The second way is to use a nested relation.
- Another way is to use the concept of joined tables.

# A Complete Set of Relational Algebra Operations

- It has been shown that the set of relational algebra operations  $\{\sigma, \pi, \cup, \rho, -, \times\}$  is a complete set; that is, any of the other original relational algebra operations can be expressed as a *sequence of operations from this set*.

# The DIVISION Operation

- An example is *Retrieve the names of employees who work on all the projects that 'John Smith' works on.*
- SMITH  $\leftarrow \sigma_{\text{Fname}='John' \text{ AND } \text{Lname}='Smith'}(\text{EMPLOYEE})$   
SMITH\_PNOS  $\leftarrow \pi_{\text{Pno}}(\text{WORKS\_ON} \bowtie_{\text{Essn}=\text{Ssn}} \text{SMITH})$   
SSN\_PNOS  $\leftarrow \pi_{\text{Essn}, \text{Pno}}(\text{WORKS\_ON})$   
SSNS(Ssn)  $\leftarrow \text{SSN\_PNOS} \div \text{SMITH\_PNOS}$   
RESULT  $\leftarrow \pi_{\text{Fname}, \text{Lname}}(\text{SSNS} * \text{EMPLOYEE})$

- In general, the DIVISION operation is applied to two relations  $R(Z)$   $\div S(X)$ , where the attributes of  $S$  are a subset of the attributes of  $R$ ; that is,  $X \subseteq Z$ .
- Let  $Y$  be the set of attributes of  $R$  that are not attributes of  $S$ ; that is,  $Y = Z - X$  (and hence  $Z = X \cup Y$ ).
- The result of DIVISION is a relation  $T(Y)$  that includes a tuple  $t$  if tuples  $t_R$  appear in  $R$  with  $t_R[Y] = t$ , and with  $t_R[X] = t_S$  for every tuple  $t_S$  in  $S$ .



- The DIVISION operation can be expressed as a sequence of  $\Pi$ ,  $\times$ , and  $-$  operations as follows:

$$T1 \leftarrow \Pi_Y(R)$$

$$T2 \leftarrow \Pi_Y((S \times T1) - R)$$

$$T \leftarrow T1 - T2$$

- Most RDBMS implementations with SQL as the primary query language do not directly implement division.