

# The Relational Algebra

Part 3

# Additional Relational Operations

# Generalized Projection

- The generalized projection operation extends the projection operation by allowing functions of attributes to be included in the projection list.
- The generalized form can be expressed as:

$$\Pi_{F_1, F_2, \dots, F_n}(R)$$

where  $F_1, F_2, \dots, F_n$  are functions over the attributes in relation  $R$  and may involve arithmetic operations and constant values.

- As an example, consider the relation

EMPLOYEE (Ssn, Salary, Deduction, Years\_service)

- A report may be required to show

Net Salary = Salary - Deduction,

Bonus = 2000 \* Years\_service, and

Tax = 0.25 \* Salary

- Then a generalized projection combined with renaming may be used as follows:

$$\text{REPORT} \leftarrow \rho_{(\text{Ssn}, \text{Net\_salary}, \text{Bonus}, \text{Tax})} (\pi_{\text{Ssn}, \text{Salary} - \text{Deduction}, 2000 * \text{Years\_service}, 0.25 * \text{Salary}} (\text{EMPLOYEE}))$$

# Aggregate Functions and Grouping

- We define an AGGREGATE FUNCTION operation, using the symbol  $\mathfrak{F}$  (pronounced *script F*), to specify types of requests as follows:

`<grouping attributes>  $\mathfrak{F}$  <function list> (R)`

- where `<grouping attributes>` is a list of attributes of the relation specified in  $R$ , and `<function list>` is a list of (`<function>` `<attribute>`) pairs.
- In each such pair, `<function>` is one of the allowed functions—such as SUM, AVERAGE, MAXIMUM, MINIMUM, COUNT—and `<attribute>` is an attribute of the relation specified by  $R$ .

- The resulting relation has the grouping attributes plus one attribute for each element in the function list.

- $\rho_{R(Dno, No\_of\_employees, Average\_sal)} (Dno \Join COUNT Ssn, AVERAGE Salary (EMPLOYEE) )$

- If no grouping attributes are specified, the functions are applied to *all the tuples* in the relation, so the resulting relation has a *single tuple only*.
- $\mathfrak{S}$  `COUNT Ssn, AVERAGE Salary (EMPLOYEE)`

# Recursive Closure Operations

- An example of a recursive operation is to retrieve all supervisees of an employee  $e$  at all levels—that is, all employees  $e'$  directly supervised by  $e$ , all employees  $e''$  directly supervised by each employee  $e'$ , all employees  $e'''$  directly supervised by each employee  $e''$ , and so on.



- It is relatively straightforward in the relational algebra to specify all employees supervised by  $e$  *at a specific level* by joining the table with itself one or more times.
- However, it is difficult to specify all supervisees at *all* levels.

# OUTER JOIN Operations

- A set of operations, called **outer joins**, were developed for the case where the user wants to keep all the tuples in  $R$ , or all those in  $S$ , or all those in both relations in the result of the JOIN, regardless of whether or not they have matching tuples in the other relation.

- The LEFT OUTER JOIN operation keeps every tuple in the *first*, or *left*, relation  $R$  in  $R \bowtie S$ ; if no matching tuple is found in  $S$ , then the attributes of  $S$  in the join result are filled or padded with NULL values.
- For example, suppose that we want a list of all employee names as well as the name of the departments they manage *if they happen to manage a department*; if they do not manage one, we can indicate it with a NULL value.
- $$\text{TEMP} \leftarrow (\text{EMPLOYEE} \bowtie_{\text{Ssn}=\text{Mgr\_ssn}} \text{DEPARTMENT})$$
$$\text{RESULT} \leftarrow \Pi_{\text{Fname}, \text{Minit}, \text{Lname}, \text{Dname}} (\text{TEMP})$$

- A similar operation, RIGHT OUTER JOIN, denoted by  $\bowtie$ , keeps every tuple in the *second*, or *right*, relation  $S$  in the result of  $R \bowtie S$ .
- A third operation, FULL OUTER JOIN, denoted by  $\bowtie$ , keeps all tuples in both the left and the right relations when no matching tuples are found, padding them with NULL values as needed.

# The OUTER UNION Operation

- The **OUTER UNION** operation was developed to take the union of tuples from two relations that have some common attributes, but are *not union (type) compatible*.

- For example, an OUTER UNION can be applied to two relations whose schemas are STUDENT (Name, Ssn, Department, Advisor) and INSTRUCTOR (Name, Ssn, Department, Rank).
- Tuples from the two relations are matched based on having the same combination of values of the shared attributes—Name, Ssn, Department.
- The resulting relation, STUDENT\_OR\_INSTRUCTOR, will have the following attributes:  
STUDENT\_OR\_INSTRUCTOR (Name, Ssn, Department, Advisor, Rank)
- Tuples appearing only in STUDENT will have a NULL for the Rank attribute, whereas tuples appearing only in INSTRUCTOR will have a NULL for the Advisor attribute.
- A tuple that exists in both relations, which represent a student who is also an instructor, will have values for all its attributes.

- Notice that the same person may still appear twice in the result.
- For example, we could have a graduate student in the Mathematics department who is an instructor in the Computer Science department.
- Although the two tuples representing that person in `STUDENT` and `INSTRUCTOR` will have the same `(Name, Ssn)` values, they will not agree on the `Department` value, and so will not be matched.
- This is because `Department` has two different meanings in `STUDENT` (the department where the person studies) and `INSTRUCTOR` (the department where the person is employed as an instructor).
- If we wanted to apply the `OUTER UNION` based on the same `(Name, Ssn)` combination only, we should rename the `Department` attribute in each table to reflect that they have different meanings and designate them as not being part of the union-compatible attributes.
- For example, we could rename the attributes as `MajorDept` in `STUDENT` and `WorkDept` in `INSTRUCTOR`.

# Examples of Queries in Relational Algebra



- **Query 1.** Retrieve the name and address of all employees who work for the 'Research' department.
- $\text{RESEARCH\_DEPT} \leftarrow \sigma_{\text{Dname} = \text{'Research'}}(\text{DEPARTMENT})$   
 $\text{RESEARCH\_EMPS} \leftarrow (\text{RESEARCH\_DEPT} \bowtie_{\text{Dnumber}=\text{Dno}} \text{EMPLOYEE})$   
 $\text{RESULT} \leftarrow \pi_{\text{Fname}, \text{Lname}, \text{Address}}(\text{RESEARCH\_EMPS})$

- **Query 2.** For every project located in 'Stafford', list the project number, the controlling department number, and the department manager's last name, address, and birth date.
- $$\begin{aligned} \text{STAFFORD\_PROJS} &\leftarrow \sigma_{\text{Plocation} = \text{' Stafford '}}(\text{PROJECT}) \\ \text{CONTR\_DEPTS} &\leftarrow (\text{STAFFORD\_PROJS} \bowtie_{\text{Dnum}=\text{Dnumber}} \text{DEPARTMENT}) \\ \text{PROJ\_DEPT\_MGRS} &\leftarrow (\text{CONTR\_DEPTS} \bowtie_{\text{Mgr\_ssn}=\text{Ssn}} \text{EMPLOYEE}) \\ \text{RESULT} &\leftarrow \pi_{\text{Pnumber, Dnum, Lname, Address, Bdate}}(\text{PROJ\_DEPT\_MGRS}) \end{aligned}$$

- **Query 3.** Find the names of employees who work on *all* the projects controlled by department number 5.
- $\text{DEPT5\_PROJS} \leftarrow \rho_{(Pno)} (\Pi_{Pnumber} (\sigma_{Dnum=5} (\text{PROJECT})))$   
 $\text{EMP\_PROJ} \leftarrow \rho_{(Ssn, Pno)} (\Pi_{Essn, Pno} (\text{WORKS\_ON}))$   
 $\text{RESULT\_EMP\_SSNS} \leftarrow \text{EMP\_PROJ} \div \text{DEPT5\_PROJS}$   
 $\text{RESULT} \leftarrow \Pi_{Lname, Fname} (\text{RESULT\_EMP\_SSNS} * \text{EMPLOYEE})$

- **Query 4.** Make a list of project numbers for projects that involve an employee whose last name is 'Smith', either as a worker or as a manager of the department that controls the project.
- $$\begin{aligned} \text{SMITHS}(\text{Essn}) &\leftarrow \Pi_{\text{Ssn}} (\sigma_{\text{Lname}='Smith'}(\text{EMPLOYEE})) \\ \text{SMITH\_WORKER\_PROJS} &\leftarrow \Pi_{\text{Pno}} (\text{WORKS\_ON} * \text{SMITHS}) \\ \text{MGRS} &\leftarrow \Pi_{\text{Lname}, \text{Dnumber}} (\text{EMPLOYEE} \bowtie_{\text{Ssn}=\text{Mgr\_ssn}} \text{DEPARTMENT}) \\ \text{SMITH\_MANAGED\_DEPTS}(\text{Dnum}) &\leftarrow \\ &\quad \Pi_{\text{Dnumber}} (\sigma_{\text{Lname}='Smith'}(\text{MGRS})) \\ \text{SMITH\_MGR\_PROJS}(\text{Pno}) &\leftarrow \Pi_{\text{Pnumber}} (\text{SMITH\_MANAGED\_DEPTS} * \text{PROJECT}) \\ \text{RESULT} &\leftarrow (\text{SMITH\_WORKER\_PROJS} \cup \text{SMITH\_MGR\_PROJS}) \end{aligned}$$

- **Query 5.** List the names of all employees with two or more dependents.
- $T1(Ssn, No\_of\_dependents) \leftarrow \text{Essn} \Join \text{COUNT } \text{Dependent\_name} (DEPENDENT)$   
 $T2 \leftarrow \sigma_{No\_of\_dependents > 2} (T1)$   
 $RESULT \leftarrow \pi_{Lname, Fname} (T2 * EMPLOYEE)$

- **Query 6.** Retrieve the names of employees who have no dependents.
- ALL\_EMPS  $\leftarrow \Pi_{Ssn} (EMPLOYEE)$   
EMPS\_WITH\_DEPS(Ssn)  $\leftarrow \Pi_{Essn} (DEPENDENT)$   
EMPS\_WITHOUT\_DEPS  $\leftarrow (ALL\_EMPS - EMPS\_WITH\_DEPS)$   
RESULT  $\leftarrow \Pi_{Lname, Fname} (EMPS\_WITHOUT\_DEPS \star EMPLOYEE)$

- **Query 7.** List the names of managers who have at least one dependent.
- $$\begin{aligned} \text{MGRS}(\text{Ssn}) &\leftarrow \Pi_{\text{Mgr\_ssn}}(\text{DEPARTMENT}) \\ \text{EMPS\_WITH\_DEPS}(\text{Ssn}) &\leftarrow \Pi_{\text{Essn}}(\text{DEPENDENT}) \\ \text{MGRS\_WITH\_DEPS} &\leftarrow (\text{MGRS} \cap \text{EMPS\_WITH\_DEPS}) \\ \text{RESULT} &\leftarrow \Pi_{\text{Lname}, \text{Fname}}(\text{MGRS\_WITH\_DEPS} * \text{EMPLOYEE}) \end{aligned}$$