

NumPy Basics: Arrays and Vectorized Computation

Part 2

The NumPy ndarray: A Multidimensional Array Object

Part 2

Data Types for ndarrays

- The *data type* or `dtype` is a special object containing the information (or metadata, data about data) the ndarray needs to interpret a chunk of memory as a particular type of data:

```
In [22]: arr1 = np.array([1, 2, 3], dtype=np.float64)
```

```
In [23]: arr1.dtype
```

```
Out[23]: dtype('float64')
```

```
In [24]: arr2 = np.array([1, 2, 3], dtype=np.int32)
```

```
In [25]: arr2.dtype
```

```
Out[25]: dtype('int32')
```

Type	Type code	Description
int8, uint8	i1, u1	Signed and unsigned 8-bit (1 byte) integer types
int16, uint16	i2, u2	Signed and unsigned 16-bit integer types
int32, uint32	i4, u4	Signed and unsigned 32-bit integer types
int64, uint64	i8, u8	Signed and unsigned 64-bit integer types

<code>float16</code>	<code>f2</code>	Half-precision floating point
<code>float32</code>	<code>f4</code> or <code>f</code>	Standard single-precision floating point; compatible with C float
<code>float64</code>	<code>f8</code> or <code>d</code>	Standard double-precision floating point; compatible with C double and Python <code>float</code> object
<code>float128</code>	<code>f16</code> or <code>g</code>	Extended-precision floating point

complex64, complex128, complex256	c8, c16, c32	Complex numbers represented by two 32, 64, or 128 floats, respectively
bool	?	Boolean type storing True and False values
object	O	Python object type; a value can be any Python object
string_	S	Fixed-length ASCII string type (1 byte per character); for example, to create a string dtype with length 10, use 'S10'
unicode_	U	Fixed-length Unicode type (number of bytes platform specific); same specification semantics as string_ (e.g., 'U10')

- You can explicitly convert or *cast* an array from one dtype to another using ndarray's `astype` method:

```
In [26]: arr = np.array([1, 2, 3, 4, 5])  
arr.dtype
```

```
Out[26]: dtype('int64')
```

```
In [27]: float_arr = arr.astype(np.float64)  
float_arr.dtype
```

```
Out[27]: dtype('float64')
```

- If I cast some floating-point numbers to be of integer dtype, the decimal part will be truncated:

```
In [28]: arr = np.array([3.7, -1.2, -2.6, 0.5, 12.9, 10.1])  
arr
```

```
Out[28]: array([ 3.7, -1.2, -2.6,  0.5, 12.9, 10.1])
```

```
In [29]: arr.astype(np.int32)
```

```
Out[29]: array([ 3, -1, -2,  0, 12, 10], dtype=int32)
```


- If you have an array of strings representing numbers, you can use `astype` to convert them to numeric form:

```
In [30]: numeric_strings = np.array(['1.25', '-9.6', '42'], dtype=np.string_)
         numeric_strings.astype(float)

Out[30]: array([ 1.25, -9.6 , 42.  ])
```

- You can also use another array's `dtype` attribute:

```
In [31]: int_array = np.arange(10)
          calibers = np.array([.22, .270, .357, .380, .44, .50], dtype=np.float64)
          int_array.astype(calibers.dtype)

Out[31]: array([0., 1., 2., 3., 4., 5., 6., 7., 8., 9.])
```

- There are shorthand type code strings you can also use to refer to a dtype:

```
In [32]: empty_uint32 = np.empty(8, dtype='u4')  
empty_uint32
```

```
Out[32]: array([      0, 1075314688,      0, 1075707904,      0,  
                1075838976,      0, 1072693248], dtype=uint32)
```

- Calling `astype` always creates a new array (a copy of the data), even if the new dtype is the same as the old dtype.

Arithmetic with NumPy Arrays

- Arrays are important because they enable you to express batch operations on data without writing any for loops.
- NumPy users call this *vectorization*.
- Any arithmetic operations between equal-size arrays applies the operation element-wise:

```
In [33]: arr = np.array([[1., 2., 3.], [4., 5., 6.]])  
arr
```

```
Out[33]: array([[1., 2., 3.],  
               [4., 5., 6.]])
```

```
In [34]: arr * arr
```

```
Out[34]: array([[ 1.,  4.,  9.],  
               [16., 25., 36.]])
```

```
In [35]: arr - arr
```

```
Out[35]: array([[0., 0., 0.],  
               [0., 0., 0.]])
```

- Arithmetic operations with scalars propagate the scalar argument to each element in the array:

```
In [36]: arr
```

```
Out[36]: array([[1., 2., 3.],  
               [4., 5., 6.]])
```

```
In [37]: 1 / arr
```

```
Out[37]: array([[1.      , 0.5      , 0.33333333],  
               [0.25     , 0.2      , 0.16666667]])
```

```
In [38]: arr ** 0.5
```

```
Out[38]: array([[1.      , 1.41421356, 1.73205081],  
               [2.      , 2.23606798, 2.44948974]])
```

- Comparisons between arrays of the same size yield boolean arrays:

```
In [39]: arr2 = np.array([[0., 4., 1.], [7., 2., 12.]])  
arr2
```

```
Out[39]: array([[ 0.,  4.,  1.],  
               [ 7.,  2., 12.]])
```

```
In [40]: arr
```

```
Out[40]: array([[1., 2., 3.],  
               [4., 5., 6.]])
```

```
In [41]: arr2 > arr
```

```
Out[41]: array([[False,  True, False],  
               [ True, False,  True]])
```

Basic Indexing and Slicing

- NumPy array indexing is a rich topic, as there are many ways you may want to select a subset of your data or individual elements.
- One-dimensional arrays are simple; on the surface they act similarly to Python lists:

```
In [42]: arr = np.arange(10)  
arr
```

```
Out[42]: array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9])
```

```
In [43]: arr[5]
```

```
Out[43]: 5
```

```
In [44]: arr[5:8]
```

```
Out[44]: array([5, 6, 7])
```

```
In [45]: arr[5:8] = 12  
arr
```

```
Out[45]: array([ 0,  1,  2,  3,  4, 12, 12, 12,  8,  9])
```


- As you can see, if you assign a scalar value to a slice, as in `arr[5:8] = 12`, the value is propagated (or *broadcasted* henceforth) to the entire selection.
- An important first distinction from Python's built-in lists is that array slices are views on the original array.
- This means that the data is not copied, and any modifications to the view will be reflected in the source array.

- To give an example of this, I first create a slice of `arr`:

```
In [46]: arr_slice = arr[5:8]  
arr_slice
```

```
Out[46]: array([12, 12, 12])
```

- Now, when I change values in `arr_slice`, the mutations are reflected in the original array `arr`:

```
In [47]: arr_slice[1] = 12345  
arr
```

```
Out[47]: array([ 0,  1,  2,  3,  4, 12, 12345, 12,  8,  
                9])
```

- The “bare” slice `[:]` will assign to all values in an array:

```
In [48]: arr_slice[:] = 64  
arr
```

```
Out[48]: array([ 0,  1,  2,  3,  4, 64, 64, 64,  8,  9])
```

- If you are new to NumPy, you might be surprised by this, especially if you have used other array programming languages that copy data more eagerly.
- As NumPy has been designed to be able to work with very large arrays, you could imagine performance and memory problems if NumPy insisted on always copying data.

- If you want a copy of a slice of an ndarray instead of a view, you will need to explicitly copy the array—for example, `arr[5:8].copy()`.

- With higher dimensional arrays, you have many more options.
- In a two-dimensional array, the elements at each index are no longer scalars but rather one-dimensional arrays:

```
In [49]: arr2d = np.array([[1, 2, 3], [4, 5, 6], [7, 8, 9]])  
arr2d[2]
```

```
Out[49]: array([7, 8, 9])
```

- Thus, individual elements can be accessed recursively.
- But that is a bit too much work, so you can pass a comma-separated list of indices to select individual elements.
- So these are equivalent:

```
In [50]: arr2d
```

```
Out[50]: array([[1, 2, 3],  
               [4, 5, 6],  
               [7, 8, 9]])
```

```
In [51]: arr2d[0][2]
```

```
Out[51]: 3
```

```
In [52]: arr2d[0, 2]
```

```
Out[52]: 3
```

- See the following figure for an illustration of indexing on a two-dimensional array.
- I find it helpful to think of axis 0 as the “rows” of the array and axis 1 as the “columns.”

		axis 1		
		0	1	2
axis 0	0	0, 0	0, 1	0, 2
	1	1, 0	1, 1	1, 2
	2	2, 0	2, 1	2, 2

- In multidimensional arrays, if you omit later indices, the returned object will be a lower dimensional ndarray consisting of all the data along the higher dimensions.
- So in the $2 \times 2 \times 3$ array `arr3d`:

```
In [53]: arr3d = np.array([[[1, 2, 3], [4, 5, 6]], [[7, 8, 9], [10, 11, 12]]])
arr3d
Out[53]: array([[[ 1,  2,  3],
                  [ 4,  5,  6]],
                [[ 7,  8,  9],
                 [10, 11, 12]]])
```

- `arr3d[0]` is a 2×3 array:

```
In [54]: arr3d[0]
Out[54]: array([[1, 2, 3],
                [4, 5, 6]])
```

- Both scalar values and arrays can be assigned to `arr3d[0]`:

```
In [55]: old_values = arr3d[0].copy()  
arr3d[0] = 42  
arr3d
```

```
Out[55]: array([[[42, 42, 42],  
                [42, 42, 42]],  
               [[ 7,  8,  9],  
                [10, 11, 12]]])
```

```
In [56]: arr3d[0] = old_values  
arr3d
```

```
Out[56]: array([[[ 1,  2,  3],  
                [ 4,  5,  6]],  
               [[ 7,  8,  9],  
                [10, 11, 12]]])
```

- Similarly, `arr3d[1, 0]` gives you all of the values whose indices start with `(1, 0)`, forming a 1-dimensional array:

```
In [57]: arr3d
Out[57]: array([[[ 1,  2,  3],
                  [ 4,  5,  6]],
                [[ 7,  8,  9],
                 [10, 11, 12]]])
```

```
In [58]: arr3d[1, 0]
Out[58]: array([7, 8, 9])
```

- This expression is the same as though we had indexed in two steps:

```
In [59]: x = arr3d[1]
          x
Out[59]: array([[ 7,  8,  9],
                 [10, 11, 12]])
```

```
In [60]: x[0]
Out[60]: array([7, 8, 9])
```

- Note that in all of these cases where subsections of the array have been selected, the returned arrays are views.