

# Data Wrangling: Join, Combine, and Reshape

Part 1

# Hierarchical Indexing

# Hierarchical Indexing

- *Hierarchical indexing* is an important feature of pandas that enables you to have multiple (two or more) index *levels* on an axis.

```
In [2]: data = pd.Series(np.random.randn(9),  
                        index=[['a', 'a', 'a', 'b', 'b', 'c', 'c', 'd', 'd'],  
                             [1, 2, 3, 1, 3, 1, 2, 2, 3]])
```

```
In [3]: data
```

```
Out[3]: a 1 -0.204708  
        2  0.478943  
        3 -0.519439  
       b 1 -0.555730  
        3  1.965781  
       c 1  1.393406  
        2  0.092908  
       d 2  0.281746  
        3  0.769023  
dtype: float64
```

- What you're seeing is a prettified view of a Series with a MultiIndex as its index.
- The “gaps” in the index display mean “use the label directly above”:

```
In [4]: data.index
```

```
Out[4]: MultiIndex(levels=[['a', 'b', 'c', 'd'], [1, 2, 3]],  
                   codes=[[0, 0, 0, 1, 1, 2, 2, 3, 3], [0, 1, 2, 0, 2, 0, 1, 1, 2]])
```

- With a hierarchically indexed object, so-called *partial* indexing is possible, enabling you to concisely select subsets of the data:

```
In [5]: data
```

```
Out[5]: a  1  -0.204708  
        2   0.478943  
        3  -0.519439  
        b  1  -0.555730  
        3   1.965781  
        c  1   1.393406  
        2   0.092908  
        d  2   0.281746  
        3   0.769023  
        dtype: float64
```

```
In [6]: data['b']
```

```
Out[6]: 1  -0.555730  
        3   1.965781  
        dtype: float64
```

```
In [7]: data['b':'c']
```

```
Out[7]: b  1  -0.555730  
        3   1.965781  
        c  1   1.393406  
        2   0.092908  
        dtype: float64
```

```
In [8]: data.loc[['b', 'd']]
```

```
Out[8]: b  1  -0.555730  
        3   1.965781  
        d  2   0.281746  
        3   0.769023  
        dtype: float64
```

- Selection is even possible from an “inner” level:

```
In [9]: data
```

```
Out[9]: a  1  -0.204708  
        2   0.478943  
        3  -0.519439  
       b  1  -0.555730  
        3   1.965781  
       c  1   1.393406  
        2   0.092908  
       d  2   0.281746  
        3   0.769023  
      dtype: float64
```

```
In [10]: data.loc[:, 2]
```

```
Out[10]: a    0.478943  
         c    0.092908  
         d    0.281746  
      dtype: float64
```

- Hierarchical indexing plays an important role in reshaping data and group-based operations like forming a pivot table.
- For example, you could rearrange the data into a DataFrame using its `unstack` method:

```
In [11]: data
```

```
Out[11]: a 1 -0.204708  
          2  0.478943  
          3 -0.519439  
        b 1 -0.555730  
          3  1.965781  
        c 1  1.393406  
          2  0.092908  
        d 2  0.281746  
          3  0.769023  
        dtype: float64
```

```
In [12]: data.unstack()
```

```
Out[12]:
```

	1	2	3
a	-0.204708	0.478943	-0.519439
b	-0.555730	NaN	1.965781
c	1.393406	0.092908	NaN
d	NaN	0.281746	0.769023

- The inverse operation of `unstack` is `stack`:

```
In [13]: data.unstack().stack()
```

```
Out[13]: a 1 -0.204708  
         2  0.478943  
         3 -0.519439  
        b 1 -0.555730  
         3  1.965781  
        c 1  1.393406  
         2  0.092908  
        d 2  0.281746  
         3  0.769023  
        dtype: float64
```



- With a DataFrame, either axis can have a hierarchical index:

```
In [14]: frame = pd.DataFrame(np.arange(12).reshape((4, 3)),  
                             index=[['a', 'a', 'b', 'b'], [1, 2, 1, 2]],  
                             columns=[['Ohio', 'Ohio', 'Colorado'],  
                                      ['Green', 'Red', 'Green']])
```

```
In [15]: frame
```

```
Out[15]:
```

		Ohio		Colorado	
		Green	Red	Green	
a	1	0	1	2	
	2	3	4	5	
b	1	6	7	8	
	2	9	10	11	

- The hierarchical levels can have names (as strings or any Python objects).
- If so, these will show up in the console output:

```
In [16]: frame.index.names = ['key1', 'key2']
```

```
In [17]: frame.columns.names = ['state', 'color']
```

```
In [18]: frame
```

```
Out[18]:
```

		state		Ohio		Colorado	
		color		Green	Red	Green	
key1	key2						
a	1		0	1		2	
	2		3	4		5	
b	1		6	7		8	
	2		9	10		11	

- With partial column indexing you can similarly select groups of columns:

In [18]:

```
frame
```

Out[18]:

		state Ohio		Colorado	
		color	Green	Red	Green
key1	key2				
a	1		0	1	2
	2		3	4	5
b	1		6	7	8
	2		9	10	11

In [19]:

```
frame['Ohio']
```

Out[19]:

		color	Green	Red
key1	key2			
a	1		0	1
	2		3	4
b	1		6	7
	2		9	10

- A `MultiIndex` can be created by itself and then reused; the columns in the preceding `DataFrame` with level names could be created like this:

```
In [21]: pd.MultiIndex.from_arrays(['Ohio', 'Ohio', 'Colorado'], ['Green', 'Red', 'Green'],  
                                   names=['state', 'color'])
```

```
Out[21]: MultiIndex(levels=[['Colorado', 'Ohio'], ['Green', 'Red']],  
                    codes=[[1, 1, 0], [0, 1, 0]],  
                    names=['state', 'color'])
```

# Reordering and Sorting Levels

- At times you will need to rearrange the order of the levels on an axis or sort the data by the values in one specific level.

- The `swaplevel` takes two level numbers or names and returns a new object with the levels interchanged (but the data is otherwise unaltered):

In [22]:

frame

Out[22]:

		state		Ohio		Colorado	
		color		Green	Red	Green	
key1	key2						
a	1			0	1		2
	2			3	4		5
b	1			6	7		8
	2			9	10		11

In [23]:

frame.swaplevel('key1', 'key2')

Out[23]:

		state		Ohio		Colorado	
		color		Green	Red	Green	
key2	key1						
1	a			0	1		2
2	a			3	4		5
1	b			6	7		8
2	b			9	10		11

- `sort_index`, on the other hand, sorts the data using only the values in a single level.

In [24]:

```
frame
```

Out[24]:

		state Ohio		Colorado	
		color Green	Red	Green	
key1	key2				
a	1	0	1	2	
	2	3	4	5	
b	1	6	7	8	
	2	9	10	11	

In [25]: `frame.sort_index(level=1)`

Out[25]:

		state Ohio		Colorado	
		color Green	Red	Green	
key1	key2				
a	1	0	1	2	
b	1	6	7	8	
a	2	3	4	5	
b	2	9	10	11	

- When swapping levels, it's not uncommon to also use `sort_index` so that the result is lexicographically sorted by the indicated level:

In [26]:

```
frame
```

Out[26]:

		state		Ohio		Colorado	
		color		Green	Red	Green	
		key1	key2				
a	1	0	1	2			
	2	3	4	5			
b	1	6	7	8			
	2	9	10	11			

In [27]:

```
frame.swaplevel(0, 1).sort_index(level=0)
```

Out[27]:

		state		Ohio		Colorado	
		color		Green	Red	Green	
		key2	key1				
1	a	0	1	2			
	b	6	7	8			
2	a	3	4	5			
	b	9	10	11			



# Summary Statistics by Level

- Many descriptive and summary statistics on DataFrame and Series have a `level` option in which you can specify the level you want to aggregate by on a particular axis.

- Consider the above DataFrame; we can aggregate by level on either the rows or columns like so:

In [28]:

```
frame
```

Out[28]:

	state	Ohio	Colorado	
	color	Green	Red	Green
key1	key2			
a	1	0	1	2
	2	3	4	5
b	1	6	7	8
	2	9	10	11

In [29]: `frame.sum(level='key2')`

Out[29]:

	state Ohio		Colorado	
	color Green	Red	Green	
key2				
1	6	8	10	
2	12	14	16	

In [30]: `frame.sum(level='color', axis=1)`

Out[30]:

		color Green Red	
		key1 key2	
a	1	2	1
	2	8	4
b	1	14	7
	2	20	10

# Indexing with a DataFrame's columns

- It's not unusual to want to use one or more columns from a DataFrame as the row index; alternatively, you may wish to move the row index into the DataFrame's columns.

- Here's an example DataFrame:

```
In [31]: frame = pd.DataFrame({'a': range(7), 'b': range(7, 0, -1),  
                               'c': ['one', 'one', 'one', 'two', 'two',  
                                     'two', 'two'],  
                               'd': [0, 1, 2, 0, 1, 2, 3]})
```

```
In [32]: frame
```

```
Out[32]:
```

	a	b	c	d
0	0	7	one	0
1	1	6	one	1
2	2	5	one	2
3	3	4	two	0
4	4	3	two	1
5	5	2	two	2
6	6	1	two	3

- DataFrame's `set_index` function will create a new DataFrame using one or more of its columns as the index:

```
In [33]: frame2 = frame.set_index(['c', 'd'])
```

```
In [34]: frame2
```

```
Out[34]:
```

		a	b
	c	d	
one	0	0	7
	1	1	6
	2	2	5
two	0	3	4
	1	4	3
	2	5	2
	3	6	1

- By default the columns are removed from the DataFrame, though you can leave them in:

```
In [35]: frame.set_index(['c', 'd'], drop=False)
```

```
Out[35]:
```

	a	b	c	d
c d				
one	0	0	7	one 0
	1	1	6	one 1
	2	2	5	one 2
two	0	3	4	two 0
	1	4	3	two 1
	2	5	2	two 2
	3	6	1	two 3

- `reset_index`, on the other hand, does the opposite of `set_index`; the hierarchical index levels are moved into the columns:

```
In [36]: frame2.reset_index()
```

```
Out[36]:
```

	c	d	a	b
0	one	0	0	7
1	one	1	1	6
2	one	2	2	5
3	two	0	3	4
4	two	1	4	3
5	two	2	5	2
6	two	3	6	1