

Data Aggregation and Group Operations

Part 3

Apply: General split-apply-combine

Part 1

- `apply` splits the object being manipulated into pieces, invokes the passed function on each piece, and then attempts to concatenate the pieces together.

- Returning to the tipping dataset from before, suppose you wanted to select the top five `tip_pct` values by group.
- First, write a function that selects the rows with the largest values in a particular column:

```
In [44]: def top(df, n=5, column='tip_pct'):
         return df.sort_values(by=column)[-n:]
```

```
In [45]: top(tips, n=6)
```

```
Out[45]:
```

	total_bill	tip	smoker	day	time	size	tip_pct
109	14.31	4.00	Yes	Sat	Dinner	2	0.279525
183	23.17	6.50	Yes	Sun	Dinner	4	0.280535
232	11.61	3.39	No	Sat	Dinner	2	0.291990
67	3.07	1.00	Yes	Sat	Dinner	1	0.325733
178	9.60	4.00	Yes	Sun	Dinner	2	0.416667
172	7.25	5.15	Yes	Sun	Dinner	2	0.710345

- Now, if we group by `smoker`, say, and call `apply` with this function, we get the following:

```
In [46]: tips.groupby('smoker').apply(top)
```

```
Out[46]:
```

		total_bill	tip	smoker	day	time	size	tip_pct
smoker								
No	88	24.71	5.85	No	Thur	Lunch	2	0.236746
	185	20.69	5.00	No	Sun	Dinner	5	0.241663
	51	10.29	2.60	No	Sun	Dinner	2	0.252672
	149	7.51	2.00	No	Thur	Lunch	2	0.266312
	232	11.61	3.39	No	Sat	Dinner	2	0.291990
Yes	109	14.31	4.00	Yes	Sat	Dinner	2	0.279525
	183	23.17	6.50	Yes	Sun	Dinner	4	0.280535
	67	3.07	1.00	Yes	Sat	Dinner	1	0.325733
	178	9.60	4.00	Yes	Sun	Dinner	2	0.416667
	172	7.25	5.15	Yes	Sun	Dinner	2	0.710345

- If you pass a function to `apply` that takes other arguments or keywords, you can pass these after the function:

```
In [47]: tips.groupby(['smoker', 'day']).apply(top, n=1, column='total_bill')
```

```
Out[47]:
```

		total_bill	tip	smoker	day	time	size	tip_pct
smoker		day						
No	Fri	94	22.75	3.25	No	Fri	Dinner	2 0.142857
	Sat	212	48.33	9.00	No	Sat	Dinner	4 0.186220
	Sun	156	48.17	5.00	No	Sun	Dinner	6 0.103799
	Thur	142	41.19	5.00	No	Thur	Lunch	5 0.121389
Yes	Fri	95	40.17	4.73	Yes	Fri	Dinner	4 0.117750
	Sat	170	50.81	10.00	Yes	Sat	Dinner	3 0.196812
	Sun	182	45.35	3.50	Yes	Sun	Dinner	3 0.077178
	Thur	197	43.11	5.00	Yes	Thur	Lunch	4 0.115982

Suppressing the Group Keys

- In the preceding examples, you see that the resulting object has a hierarchical index formed from the group keys along with the indexes of each piece of the original object.
- You can disable this by passing `group_keys=False` to `groupby`:

```
In [48]: tips.groupby('smoker', group_keys=False).apply(top)
```

```
Out[48]:
```

	total_bill	tip	smoker	day	time	size	tip_pct
88	24.71	5.85	No	Thur	Lunch	2	0.236746
185	20.69	5.00	No	Sun	Dinner	5	0.241663
51	10.29	2.60	No	Sun	Dinner	2	0.252672
149	7.51	2.00	No	Thur	Lunch	2	0.266312
232	11.61	3.39	No	Sat	Dinner	2	0.291990
109	14.31	4.00	Yes	Sat	Dinner	2	0.279525
183	23.17	6.50	Yes	Sun	Dinner	4	0.280535
67	3.07	1.00	Yes	Sat	Dinner	1	0.325733
178	9.60	4.00	Yes	Sun	Dinner	2	0.416667
172	7.25	5.15	Yes	Sun	Dinner	2	0.710345

Quantile and Bucket Analysis

- pandas has some tools, in particular `cut` and `qcut`, for slicing data up into buckets with bins of your choosing or by sample quantiles.
- Combining these functions with `groupby` makes it convenient to perform bucket or quantile analysis on a dataset.

- Consider a simple random dataset and an equal-length bucket categorization using `cut`:

```
In [49]: frame = pd.DataFrame({'data1': np.random.randn(1000),  
                               'data2': np.random.randn(1000)})  
quartiles = pd.cut(frame.data1, 4)  
quartiles[:10]
```

```
Out[49]: 0    (-1.23, 0.489]  
1    (-2.956, -1.23]  
2    (-1.23, 0.489]  
3    (0.489, 2.208]  
4    (-1.23, 0.489]  
5    (0.489, 2.208]  
6    (-1.23, 0.489]  
7    (-1.23, 0.489]  
8    (0.489, 2.208]  
9    (0.489, 2.208]  
Name: data1, dtype: category  
Categories (4, interval[float64]): [(-2.956, -1.23] < (-1.23, 0.489] < (0.489, 2.208] < (2.208, 3.928]]
```

- The `Categorical` object returned by `cut` can be passed directly to `groupby`.
- So we could compute a set of statistics for the `data2` column like so:

```
In [50]: def get_stats(group):  
         return {'min': group.min(), 'max': group.max(),  
                 'count': group.count(), 'mean': group.mean()}
```

```
In [51]: grouped = frame.data2.groupby(quartiles)
```

```
In [52]: grouped.apply(get_stats).unstack()
```

Out[52]:

	count	max	mean	min
data1				
(-2.956, -1.23]	95.0	1.670835	-0.039521	-3.399312
(-1.23, 0.489]	598.0	3.260383	-0.002051	-2.989741
(0.489, 2.208]	297.0	2.954439	0.081822	-3.745356
(2.208, 3.928]	10.0	1.765640	0.024750	-1.929776

- These were equal-length buckets; to compute equal-size buckets based on sample quantiles, use `qcut`.
- We'll pass `labels=False` to just get quantile numbers:

```
In [53]: # Return quantile numbers
grouping = pd.qcut(frame.data1, 10, labels=False)
grouped = frame.data2.groupby(grouping)
grouped.apply(get_stats).unstack()
```

Out[53]:

	count	max	mean	min
data1				
0	100.0	1.670835	-0.049902	-3.399312
1	100.0	2.628441	0.030989	-1.950098
2	100.0	2.527939	-0.067179	-2.925113
3	100.0	3.260383	0.065713	-2.315555
4	100.0	2.074345	-0.111653	-2.047939
5	100.0	2.184810	0.052130	-2.989741
6	100.0	2.458842	-0.021489	-2.223506
7	100.0	2.954439	-0.026459	-3.056990
8	100.0	2.735527	0.103406	-3.745356
9	100.0	2.377020	0.220122	-2.064111

Example: Filling Missing Values with Group-Specific Values

- When cleaning up missing data, in some cases you will replace data observations using `dropna`, but in others you may want to impute (fill in) the null (NA) values using a fixed value or some value derived from the data.

```
In [54]: s = pd.Series(np.random.randn(6))  
s[::2] = np.nan  
s
```

```
Out[54]: 0      NaN  
1   -0.125921  
2      NaN  
3   -0.884475  
4      NaN  
5    0.227290  
dtype: float64
```

```
In [55]: s.fillna(s.mean())
```

```
Out[55]: 0   -0.261035  
1   -0.125921  
2   -0.261035  
3   -0.884475  
4   -0.261035  
5    0.227290  
dtype: float64
```

- Suppose you need the fill value to vary by group.
- One way to do this is to group the data and use `apply` with a function that calls `fillna` on each data chunk.
- Here is some sample data on US states divided into eastern and western regions:

```
In [56]: states = ['Ohio', 'New York', 'Vermont', 'Florida',  
                  'Oregon', 'Nevada', 'California', 'Idaho']  
group_key = ['East'] * 4 + ['West'] * 4  
data = pd.Series(np.random.randn(8), index=states)  
data
```

```
Out[56]: Ohio          0.922264  
New York       -2.153545  
Vermont        -0.365757  
Florida        -0.375842  
Oregon          0.329939  
Nevada          0.981994  
California      1.105913  
Idaho          -1.613716  
dtype: float64
```

- Let's set some values in the data to be missing:

```
In [57]: data[['Vermont', 'Nevada', 'Idaho']] = np.nan  
data
```

```
Out[57]: Ohio      0.922264  
New York -2.153545  
Vermont    NaN  
Florida  -0.375842  
Oregon    0.329939  
Nevada    NaN  
California 1.105913  
Idaho     NaN  
dtype: float64
```

```
In [58]: data.groupby(group_key).mean()
```

```
Out[58]: East    -0.535707  
West      0.717926  
dtype: float64
```

- We can fill the NA values using the group means like so:

```
In [59]: fill_mean = lambda g: g.fillna(g.mean())  
data.groupby(group_key).apply(fill_mean)
```

```
Out[59]: Ohio          0.922264  
New York    -2.153545  
Vermont     -0.535707  
Florida     -0.375842  
Oregon       0.329939  
Nevada       0.717926  
California   1.105913  
Idaho        0.717926  
dtype: float64
```

- In another case, you might have predefined fill values in your code that vary by group.
- Since the groups have a `name` attribute set internally, we can use that:

```
In [60]: fill_values = {'East': 0.5, 'West': -1}
         fill_func = lambda g: g.fillna(fill_values[g.name])
         data.groupby(group_key).apply(fill_func)
```

```
Out[60]: Ohio          0.922264
         New York      -2.153545
         Vermont        0.500000
         Florida       -0.375842
         Oregon         0.329939
         Nevada        -1.000000
         California     1.105913
         Idaho         -1.000000
         dtype: float64
```


Example: Random Sampling and Permutation

- Suppose you wanted to draw a random sample (with or without replacement) from a large dataset for Monte Carlo simulation purposes or some other application.
- There are a number of ways to perform the “draws”; here we use the `sample` method for Series.

- To demonstrate, here's a way to construct a deck of English-style playing cards:

```
In [61]: # Hearts, Spades, Clubs, Diamonds
suits = ['H', 'S', 'C', 'D']
card_val = (list(range(1, 11)) + [10] * 3) * 4
base_names = ['A'] + list(range(2, 11)) + ['J', 'K', 'Q']
cards = []
for suit in ['H', 'S', 'C', 'D']:
    cards.extend(str(num) + suit for num in base_names)

deck = pd.Series(card_val, index=cards)
```

- So now we have a Series of length 52 whose index contains card names and values are the ones used in Blackjack and other games (to keep things simple, we just let the ace 'A' be 1):

```
In [62]: deck[:13]
```

```
Out[62]: AH      1  
         2H      2  
         3H      3  
         4H      4  
         5H      5  
         6H      6  
         7H      7  
         8H      8  
         9H      9  
        10H     10  
         JH     10  
         KH     10  
         QH     10  
        dtype: int64
```

- Drawing a hand of five cards from the deck could be written as:

```
In [63]: def draw(deck, n=5):  
         return deck.sample(n)
```

```
In [64]: draw(deck)
```

```
Out[64]: AD      1  
         8C      8  
         5H      5  
         KC     10  
         2C      2  
         dtype: int64
```

- Suppose you wanted two random cards from each suit.
- Because the suit is the last character of each card name, we can group based on this and use `apply`:

```
In [65]: get_suit = lambda card: card[-1] # last letter is suit
```

```
In [66]: deck.groupby(get_suit).apply(draw, n=2)
```

```
Out[66]: C  2C    2  
          3C    3  
D   KD    10  
      8D    8  
H   KH    10  
      3H    3  
S   2S    2  
      4S    4  
dtype: int64
```

- Alternatively, we could write:

```
In [67]: deck.groupby(get_suit, group_keys=False).apply(draw, n=2)
```

```
Out[67]: KC      10  
         JC      10  
         AD       1  
         5D       5  
         5H       5  
         6H       6  
         7S       7  
         KS      10  
         dtype: int64
```