

Data Aggregation and Group Operations

Part 4

Apply: General split-apply-combine

Part 2

Example: Group Weighted Average and Correlation

- Under the split-apply-combine paradigm of `groupby`, operations between columns in a DataFrame or two Series, such as a group weighted average, are possible.

```
In [68]: df = pd.DataFrame({'category': ['a', 'a', 'a', 'a',  
                                         'b', 'b', 'b', 'b'],  
                           'data': np.random.randn(8),  
                           'weights': np.random.rand(8)})  
df
```

Out[68]:

	category	data	weights
0	a	1.561587	0.957515
1	a	1.219984	0.347267
2	a	-0.482239	0.581362
3	a	0.315667	0.217091
4	b	-0.047852	0.894406
5	b	-0.454145	0.918564
6	b	-0.556774	0.277825
7	b	0.253321	0.955905

- The group weighted average by `category` would then be:

```
In [69]: grouped = df.groupby('category')  
get_wavg = lambda g: np.average(g['data'], weights=g['weights'])  
grouped.apply(get_wavg)
```

```
Out[69]: category  
a      0.811643  
b     -0.122262  
dtype: float64
```

- As another example, consider a financial dataset originally obtained from Yahoo! Finance containing end-of-day prices for a few stocks and the S&P 500 index (the SPX symbol):

```
In [70]: close_px = pd.read_csv('examples/stock_px_2.csv', parse_dates=True,  
                                index_col=0)
```

```
In [71]: close_px.info()  
  
<class 'pandas.core.frame.DataFrame'>  
DatetimeIndex: 2214 entries, 2003-01-02 to 2011-10-14  
Data columns (total 4 columns):  
AAPL      2214 non-null float64  
MSFT      2214 non-null float64  
XOM        2214 non-null float64  
SPX        2214 non-null float64  
dtypes: float64(4)  
memory usage: 86.5 KB
```

```
In [72]: close_px[-4:]
```

Out[72]:

	AAPL	MSFT	XOM	SPX
2011-10-11	400.29	27.00	76.27	1195.54
2011-10-12	402.19	26.96	77.16	1207.25
2011-10-13	408.43	27.18	76.37	1203.66
2011-10-14	422.00	27.27	78.11	1224.58

- One task of interest might be to compute a DataFrame consisting of the yearly correlations of daily returns (computed from percent changes) with SPX.
- As one way to do this, we first create a function that computes the pairwise correlation of each column with the 'SPX' column:

```
In [73]: spx_corr = lambda x: x.corrwith(x['SPX'])
```

- Next, we compute percent change on `close_px` using `pct_change`:

```
In [74]: rets = close_px.pct_change().dropna()
```

- Lastly, we group these percent changes by year, which can be extracted from each row label with a one-line function that returns the `year` attribute of each `datetime` label:

```
In [75]: get_year = lambda x: x.year  
by_year = rets.groupby(get_year)  
by_year.apply(spx_corr)
```

Out[75]:

	AAPL	MSFT	XOM	SPX
2003	0.541124	0.745174	0.661265	1.0
2004	0.374283	0.588531	0.557742	1.0
2005	0.467540	0.562374	0.631010	1.0
2006	0.428267	0.406126	0.518514	1.0
2007	0.508118	0.658770	0.786264	1.0
2008	0.681434	0.804626	0.828303	1.0
2009	0.707103	0.654902	0.797921	1.0
2010	0.710105	0.730118	0.839057	1.0
2011	0.691931	0.800996	0.859975	1.0

- You could also compute inter-column correlations.
- Here we compute the annual correlation between Apple and Microsoft:

```
In [76]: by_year.apply(lambda g: g['AAPL'].corr(g['MSFT']))
```

```
Out[76]: 2003    0.480868  
         2004    0.259024  
         2005    0.300093  
         2006    0.161735  
         2007    0.417738  
         2008    0.611901  
         2009    0.432738  
         2010    0.571946  
         2011    0.581987  
         dtype: float64
```

Example: Group-Wise Linear Regression

- In the same theme as the previous example, you can use `groupby` to perform more complex group-wise statistical analysis, as long as the function returns a pandas object or scalar value.
- For example, we can define the following `regress` function (using the `statsmodels` econometrics library), which executes an ordinary least squares (OLS) regression on each chunk of data:

```
In [77]: import statsmodels.api as sm
def regress(data, yvar, xvars):
    Y = data[yvar]
    X = data[xvars]
    X['intercept'] = 1.
    result = sm.OLS(Y, X).fit()
    return result.params
```

- Now, to run a yearly linear regression of AAPL on SPX returns, execute:

```
In [78]: by_year.apply(regress, 'AAPL', ['SPX'])
```

```
Out[78]:
```

	SPX	intercept
2003	1.195406	0.000710
2004	1.363463	0.004201
2005	1.766415	0.003246
2006	1.645496	0.000080
2007	1.198761	0.003438
2008	0.968016	-0.001110
2009	0.879103	0.002954
2010	1.052608	0.001261
2011	0.806605	0.001514

Pivot Tables and Cross-Tabulation

- A *pivot table* is a data summarization tool frequently found in spreadsheet programs and other data analysis software.
- It aggregates a table of data by one or more keys, arranging the data in a rectangle with some of the group keys along the rows and some along the columns.
- Pivot tables in Python with pandas are made possible through the `groupby` facility combined with reshape operations utilizing hierarchical indexing.
- `DataFrame` has a `pivot_table` method, and there is also a top-level `pandas.pivot_table` function.
- In addition to providing a convenience interface to `groupby`, `pivot_table` can add partial totals, also known as *margins*.

- Returning to the tipping dataset, suppose you wanted to compute a table of group means (the default `pivot_table` aggregation type) arranged by `day` and `smoker` on the rows:

```
In [79]: tips.pivot_table(index=['day', 'smoker'])
```

```
Out[79]:
```

		size	tip	tip_pct	total_bill
day	smoker				
Fri	No	2.250000	2.812500	0.151650	18.420000
	Yes	2.066667	2.714000	0.174783	16.813333
Sat	No	2.555556	3.102889	0.158048	19.661778
	Yes	2.476190	2.875476	0.147906	21.276667
Sun	No	2.929825	3.167895	0.160113	20.506667
	Yes	2.578947	3.516842	0.187250	24.120000
Thur	No	2.488889	2.673778	0.160298	17.113111
	Yes	2.352941	3.030000	0.163863	19.190588

- Now, suppose we want to aggregate only `tip_pct` and `size`, and additionally group by `time`.
- We'll put `smoker` in the table columns and `day` in the rows:

```
In [80]: tips.pivot_table(['tip_pct', 'size'], index=['time', 'day'],
                          columns='smoker')
```

Out[80]:

		size		tip_pct	
		smoker		No	Yes
time	day				
Dinner	Fri		2.000000	2.222222	0.139622 0.165347
	Sat		2.555556	2.476190	0.158048 0.147906
	Sun		2.929825	2.578947	0.160113 0.187250
	Thur		2.000000	NaN	0.159744 NaN
Lunch	Fri		3.000000	1.833333	0.187735 0.188937
	Thur		2.500000	2.352941	0.160311 0.163863

- We could augment this table to include partial totals by passing `margins=True`.
- This has the effect of adding `All` row and column labels, with corresponding values being the group statistics for all the data within a single tier:

```
In [81]: tips.pivot_table(['tip_pct', 'size'], index=['time', 'day'],
                          columns='smoker', margins=True)
```

Out[81]:

		size			tip_pct		
	smoker	No	Yes	All	No	Yes	All
time	day						
Dinner	Fri	2.000000	2.222222	2.166667	0.139622	0.165347	0.158916
	Sat	2.555556	2.476190	2.517241	0.158048	0.147906	0.153152
	Sun	2.929825	2.578947	2.842105	0.160113	0.187250	0.166897
	Thur	2.000000	NaN	2.000000	0.159744	NaN	0.159744
Lunch	Fri	3.000000	1.833333	2.000000	0.187735	0.188937	0.188765
	Thur	2.500000	2.352941	2.459016	0.160311	0.163863	0.161301
All		2.668874	2.408602	2.569672	0.159328	0.163196	0.160803

- To use a different aggregation function, pass it to `aggfunc`.
- For example, `'count'` or `len` will give you a cross-tabulation (count or frequency) of group sizes:

```
In [82]: tips.pivot_table('tip_pct', index=['time', 'smoker'], columns='day',  
                        aggfunc=len, margins=True)
```

Out[82]:

		day				
		Fri	Sat	Sun	Thur	All
time	smoker					
Dinner	No	3.0	45.0	57.0	1.0	106.0
	Yes	9.0	42.0	19.0	NaN	70.0
Lunch	No	1.0	NaN	NaN	44.0	45.0
	Yes	6.0	NaN	NaN	17.0	23.0
All		19.0	87.0	76.0	62.0	244.0

- If some combinations are empty (or otherwise NA), you may wish to pass a `fill_value`:

```
In [83]: tips.pivot_table('tip_pct', index=['time', 'size', 'smoker'],
                          columns='day', aggfunc='mean', fill_value=0)
```

Out[83]:

		day	Fri	Sat	Sun	Thur
time	size	smoker				
Dinner	1	No	0.000000	0.137931	0.000000	0.000000
		Yes	0.000000	0.325733	0.000000	0.000000
	2	No	0.139622	0.162705	0.168859	0.159744
		Yes	0.171297	0.148668	0.207893	0.000000
	3	No	0.000000	0.154661	0.152663	0.000000
		Yes	0.000000	0.144995	0.152660	0.000000
	4	No	0.000000	0.150096	0.148143	0.000000
		Yes	0.117750	0.124515	0.193370	0.000000
	5	No	0.000000	0.000000	0.206928	0.000000
		Yes	0.000000	0.106572	0.065660	0.000000
...
Lunch	1	No	0.000000	0.000000	0.000000	0.181728
		Yes	0.223776	0.000000	0.000000	0.000000
	2	No	0.000000	0.000000	0.000000	0.166005
		Yes	0.181969	0.000000	0.000000	0.158843
	3	No	0.187735	0.000000	0.000000	0.084246
		Yes	0.000000	0.000000	0.000000	0.204952
	4	No	0.000000	0.000000	0.000000	0.138919
		Yes	0.000000	0.000000	0.000000	0.155410
	5	No	0.000000	0.000000	0.000000	0.121389
	6	No	0.000000	0.000000	0.000000	0.173706

21 rows × 4 columns

Cross-Tabulations: Crosstab

- A cross-tabulation (or `crosstab` for short) is a special case of a pivot table that computes group frequencies.
- Here is an example:

In [85]:

```
data
```

Out[85]:

	Sample	Nationality	Handedness
0	1	USA	Right-handed
1	2	Japan	Left-handed
2	3	USA	Right-handed
3	4	Japan	Right-handed
4	5	Japan	Left-handed
5	6	Japan	Right-handed
6	7	USA	Right-handed
7	8	USA	Left-handed
8	9	Japan	Right-handed
9	10	USA	Right-handed

- As part of some survey analysis, we might want to summarize this data by nationality and handedness.
- You could use `pivot_table` to do this, but the `pandas.crosstab` function can be more convenient:

```
In [86]: pd.crosstab(data.Nationality, data.Handedness, margins=True)
```

```
Out[86]:
```

Handedness	Left-handed	Right-handed	All
Nationality			
Japan	2	3	5
USA	1	4	5
All	3	7	10

- The first two arguments to `crosstab` can each either be an array or Series or a list of arrays.
- As in the tips data:

```
In [87]: pd.crosstab([tips.time, tips.day], tips.smoker, margins=True)
```

```
Out[87]:
```

		smoker		
		No	Yes	All
time	day			
Dinner	Fri	3	9	12
	Sat	45	42	87
	Sun	57	19	76
	Thur	1	0	1
Lunch	Fri	1	6	7
	Thur	44	17	61
All		151	93	244